

Backtesting and error metrics for modern time series forecasting

PyData London 2024
June 2024
Kishan Manani

About me

- Machine learning and data science consultant
- PhD Physics: Modelling and large scale time series analysis abnormal heart rhythms
- Online course developer, see courses:
trainindata.com/p/forecasting-specialization



Kishan Manani, PhD

 KishManani

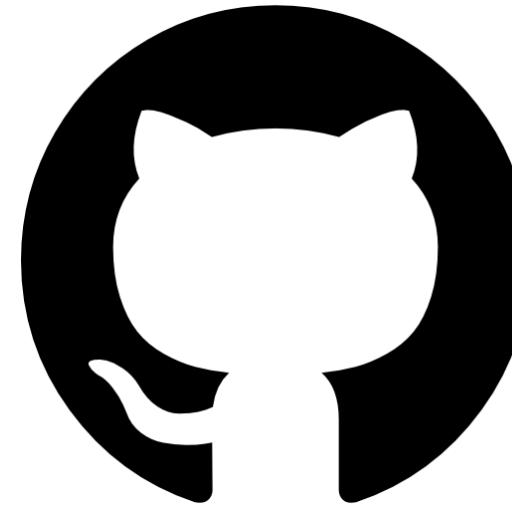
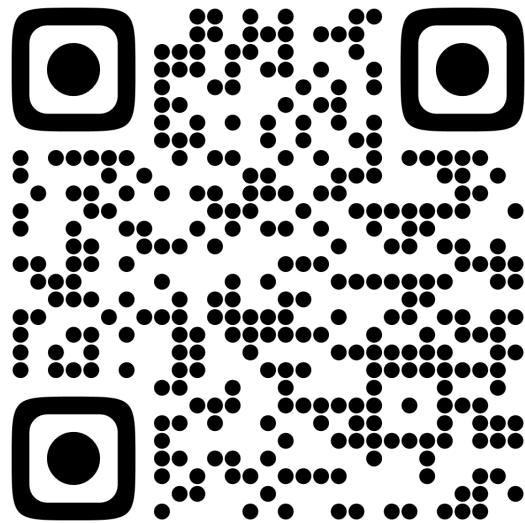
 @KishManani

 In/kishanmanani

 medium.com/@kish.manani

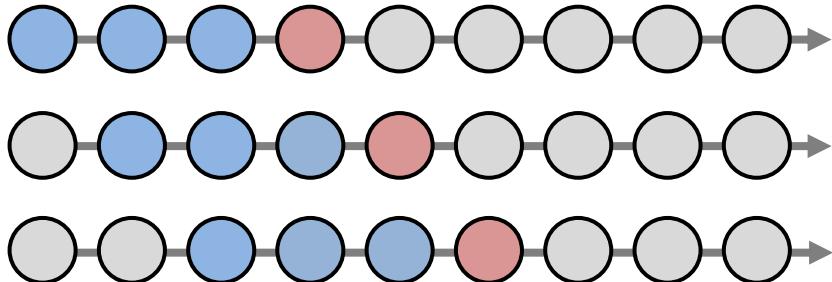
Slides

<https://github.com/KishManani/PyDataLondon2024>



What is this talk about?

Backtesting



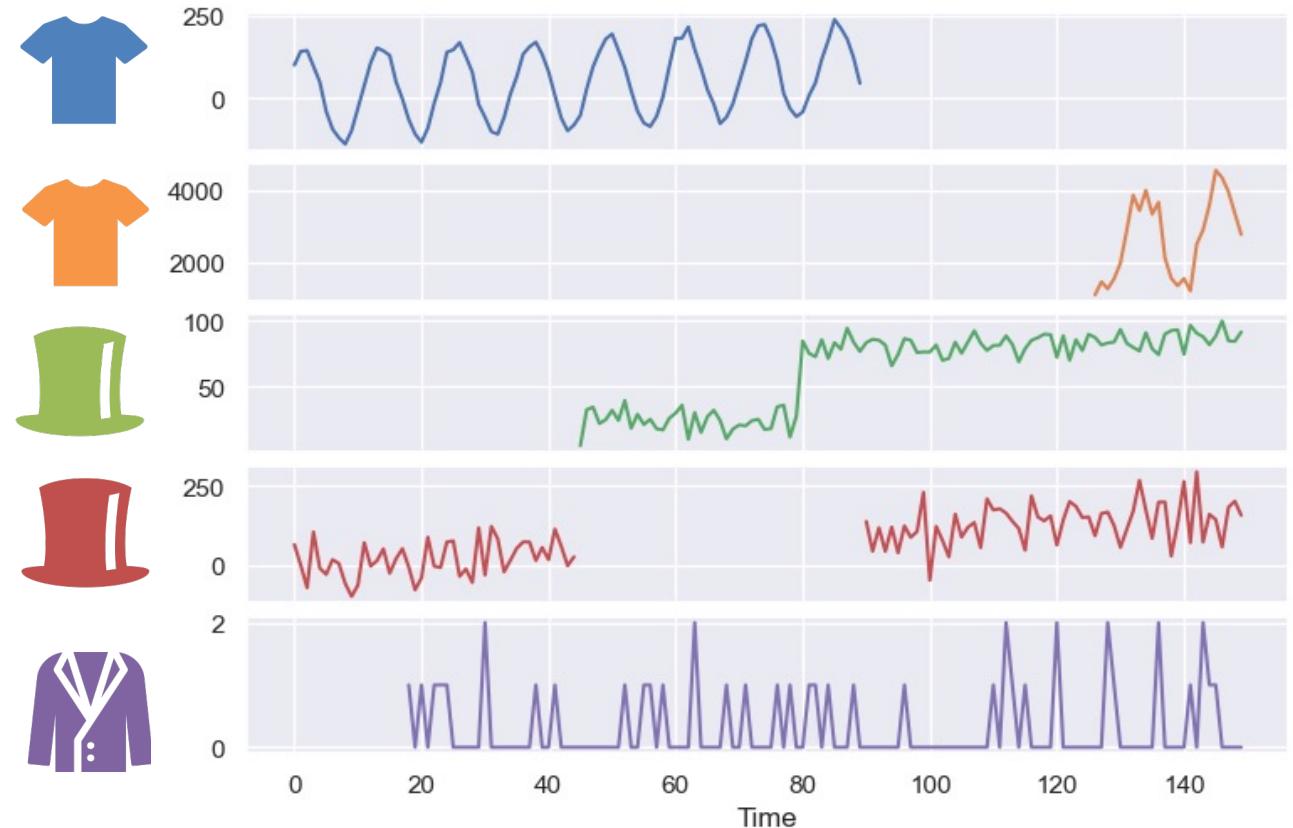
Error metrics

$$MAE = \frac{1}{N} \sum_i |e_i|$$



$$WAPE = \frac{1}{\sum_t w_t} \sum_t w_t |p_t|$$

Number of products sold



Contents



Modern time series forecasting



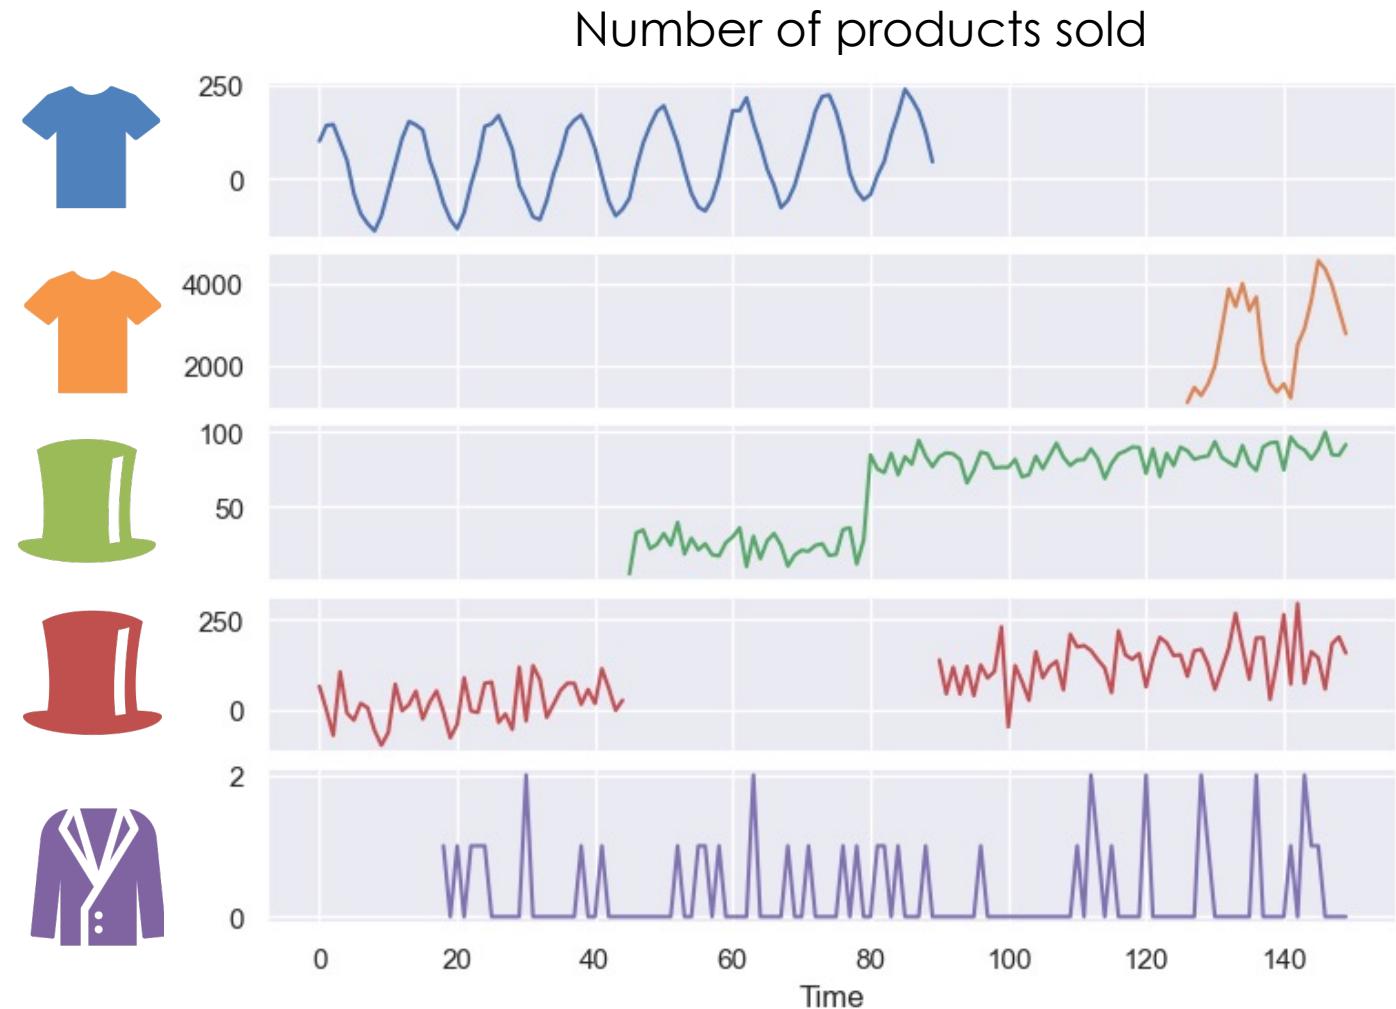
Backtesting



Error metrics

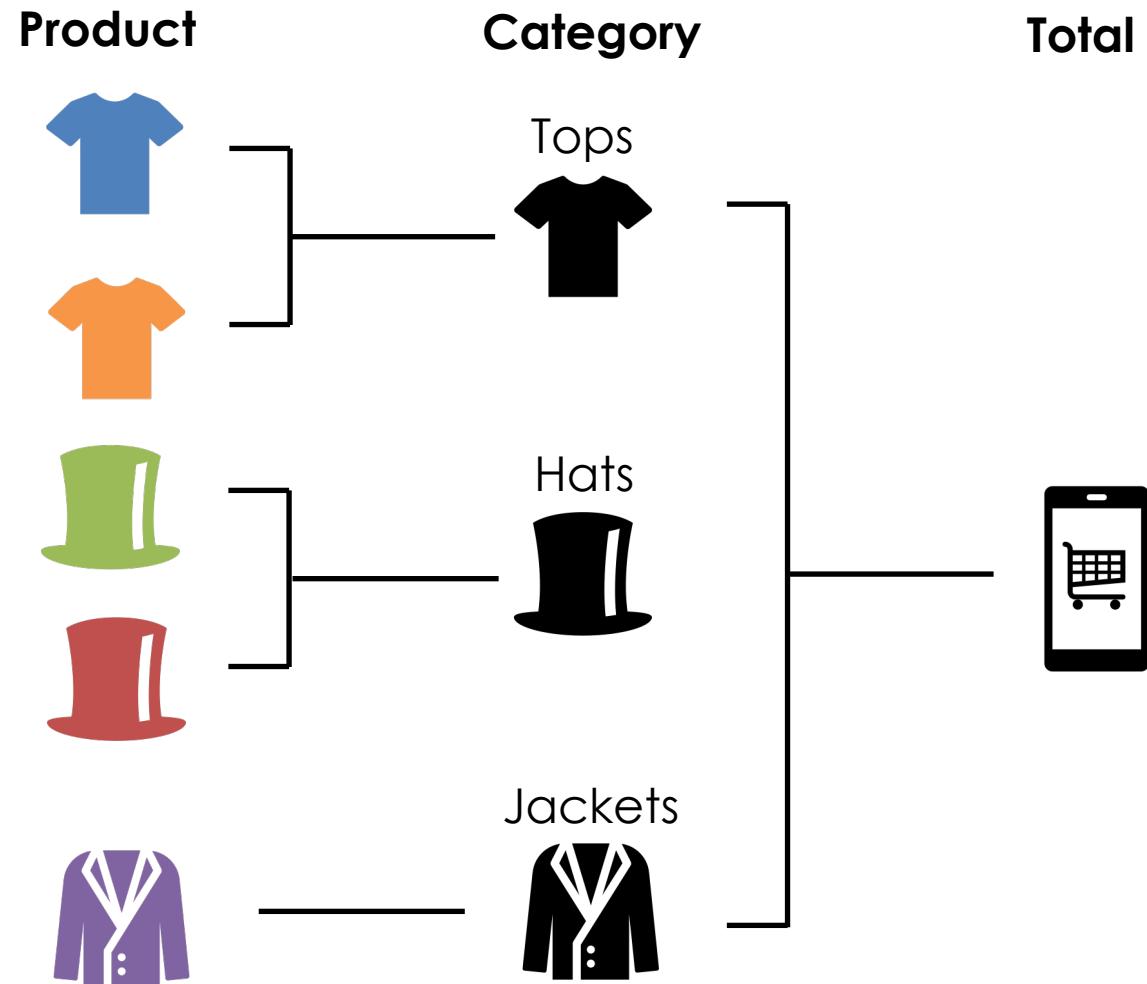
Modern time series forecasting

- Large number of related time series
- Multiple characteristics
 - Trend
 - Seasonality
 - Outliers
 - Intermittency
 - Missing data
 - Hierarchical



Modern time series forecasting

- Large number of related time series
- Multiple characteristics
 - Trend
 - Seasonality
 - Outliers
 - Intermittency
 - Missing data
 - Hierarchical

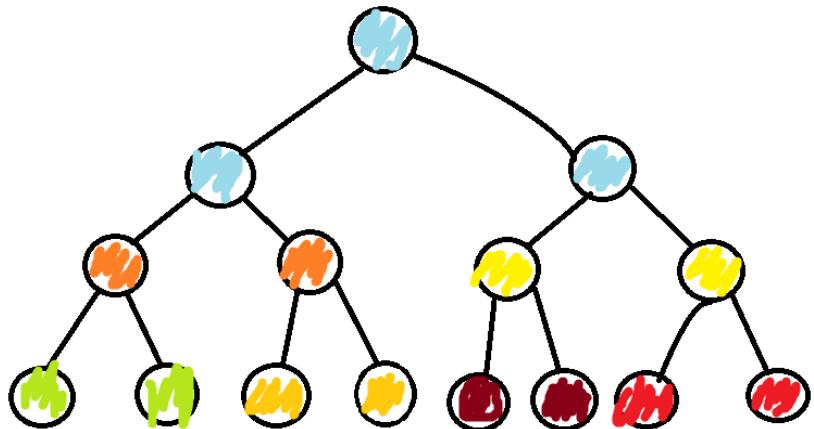


Modern time series forecasting

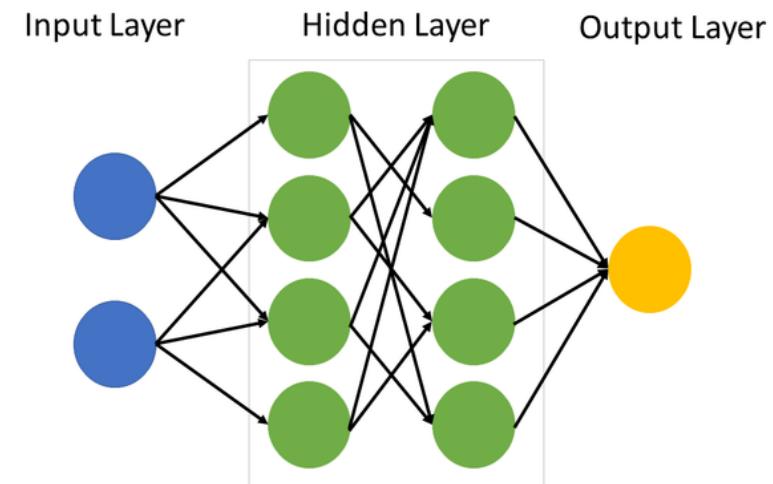


Time	Product ID	Sales
...
1	1	10
2	1	12
3	1	343
4	1	54
...
140	2	545
141	2	53
142	2	32
...

Modern time series forecasting

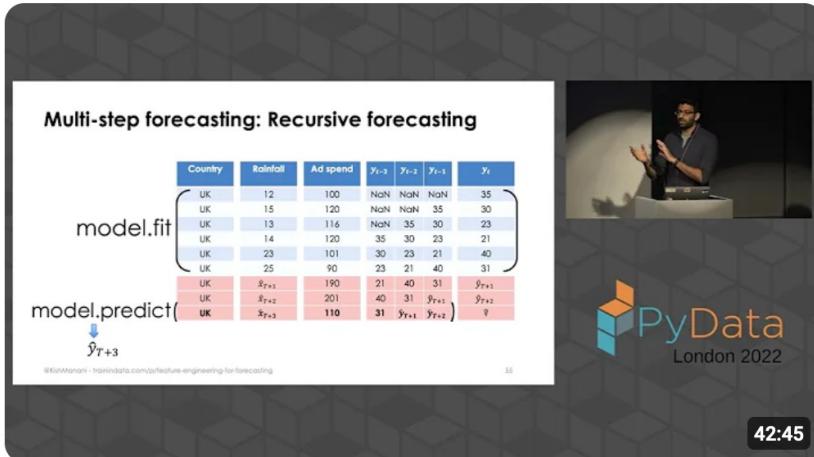


Traditional ML models



Deep learning

Modern time series forecasting



Kishan Manani - Feature Engineering for Time Series Forecasting | PyData London 2022

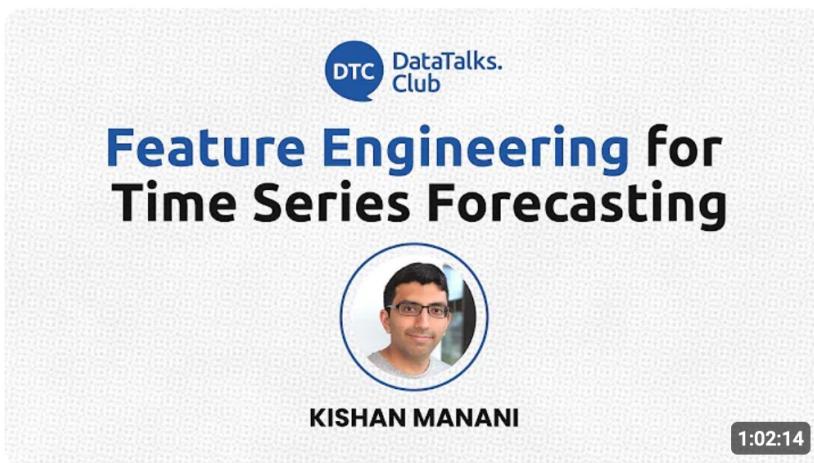
71K views • 1 year ago



Kishan Manani present: Feature Engineering for Time Series Forecasting To use our favourite supervised learning models...



Intro | About this talk | Why use machine learning for forecasting? | Don't neglect simpl... 22 chapters ▾



Feature Engineering for Time Series Forecasting - Kishan Manani

23K views • Streamed 1 year ago



In this podcast episode, we talked with Kishan Manani about feature engineering for time series forecasting. 0:00 Introdu...



Introduction and Welcome | Speaker Introduction | Topic Introduction: Feature...

19 chapters ▾

Contents



Modern time series forecasting



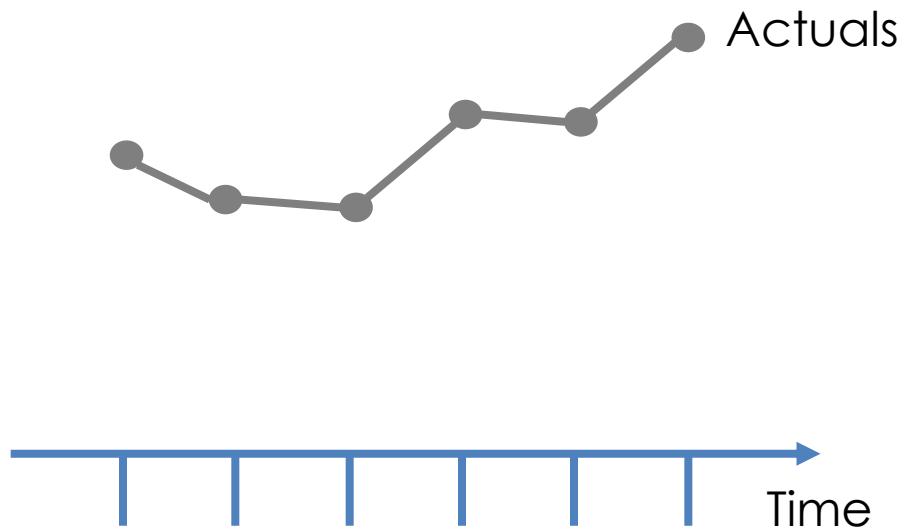
Backtesting



Error metrics

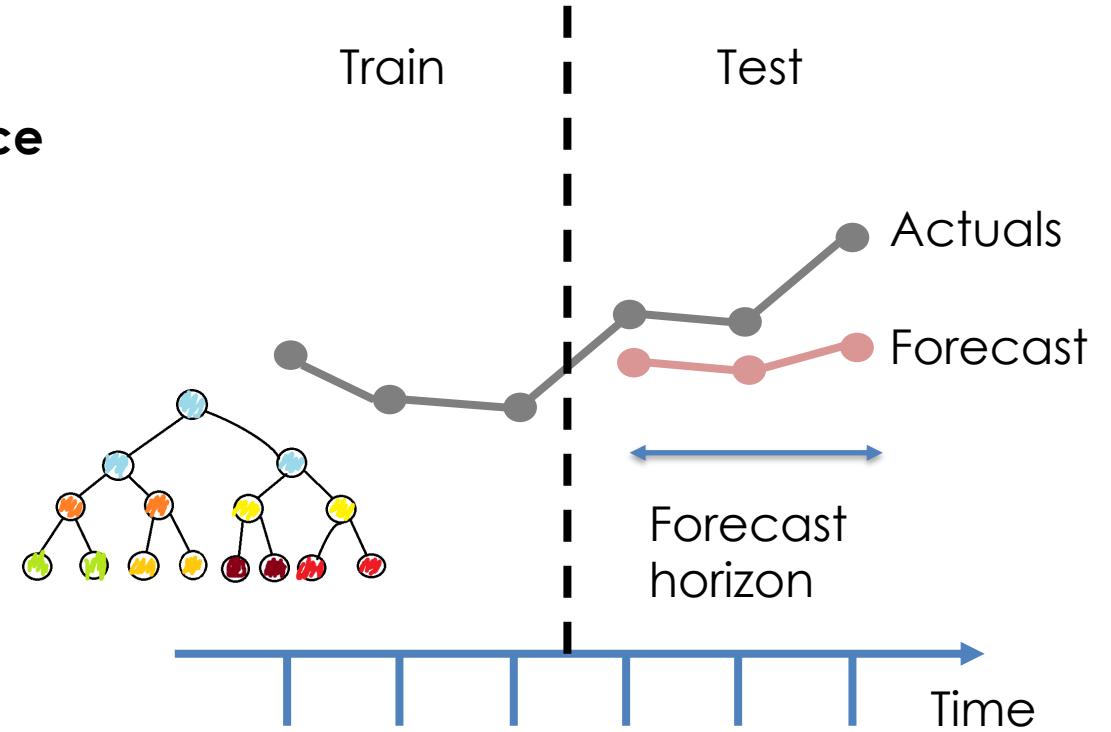
Backtesting

- Backtesting is where we **measure the performance** of a **forecasting model** on **historic data**.
- We want to **estimate the performance** of the model **on future data**.



Backtesting

- Backtesting is where we **measure the performance** of a **forecasting model** on **historic data**.
- We want to **estimate the performance** of the model **on future data**.



Backtesting

- Backtesting is where we **measure the performance** of a **forecasting model** on **historic data**.
- We want to **estimate the performance** of the model **on future data**.
- It is common to randomly shuffle data into train and test sets.

Train 

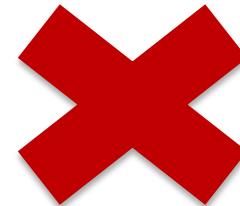
Test 

x1	x2	y
...	...	0
...	...	1
...	...	0
...	...	0
...	...	0
...	...	1
...	...	0

Backtesting

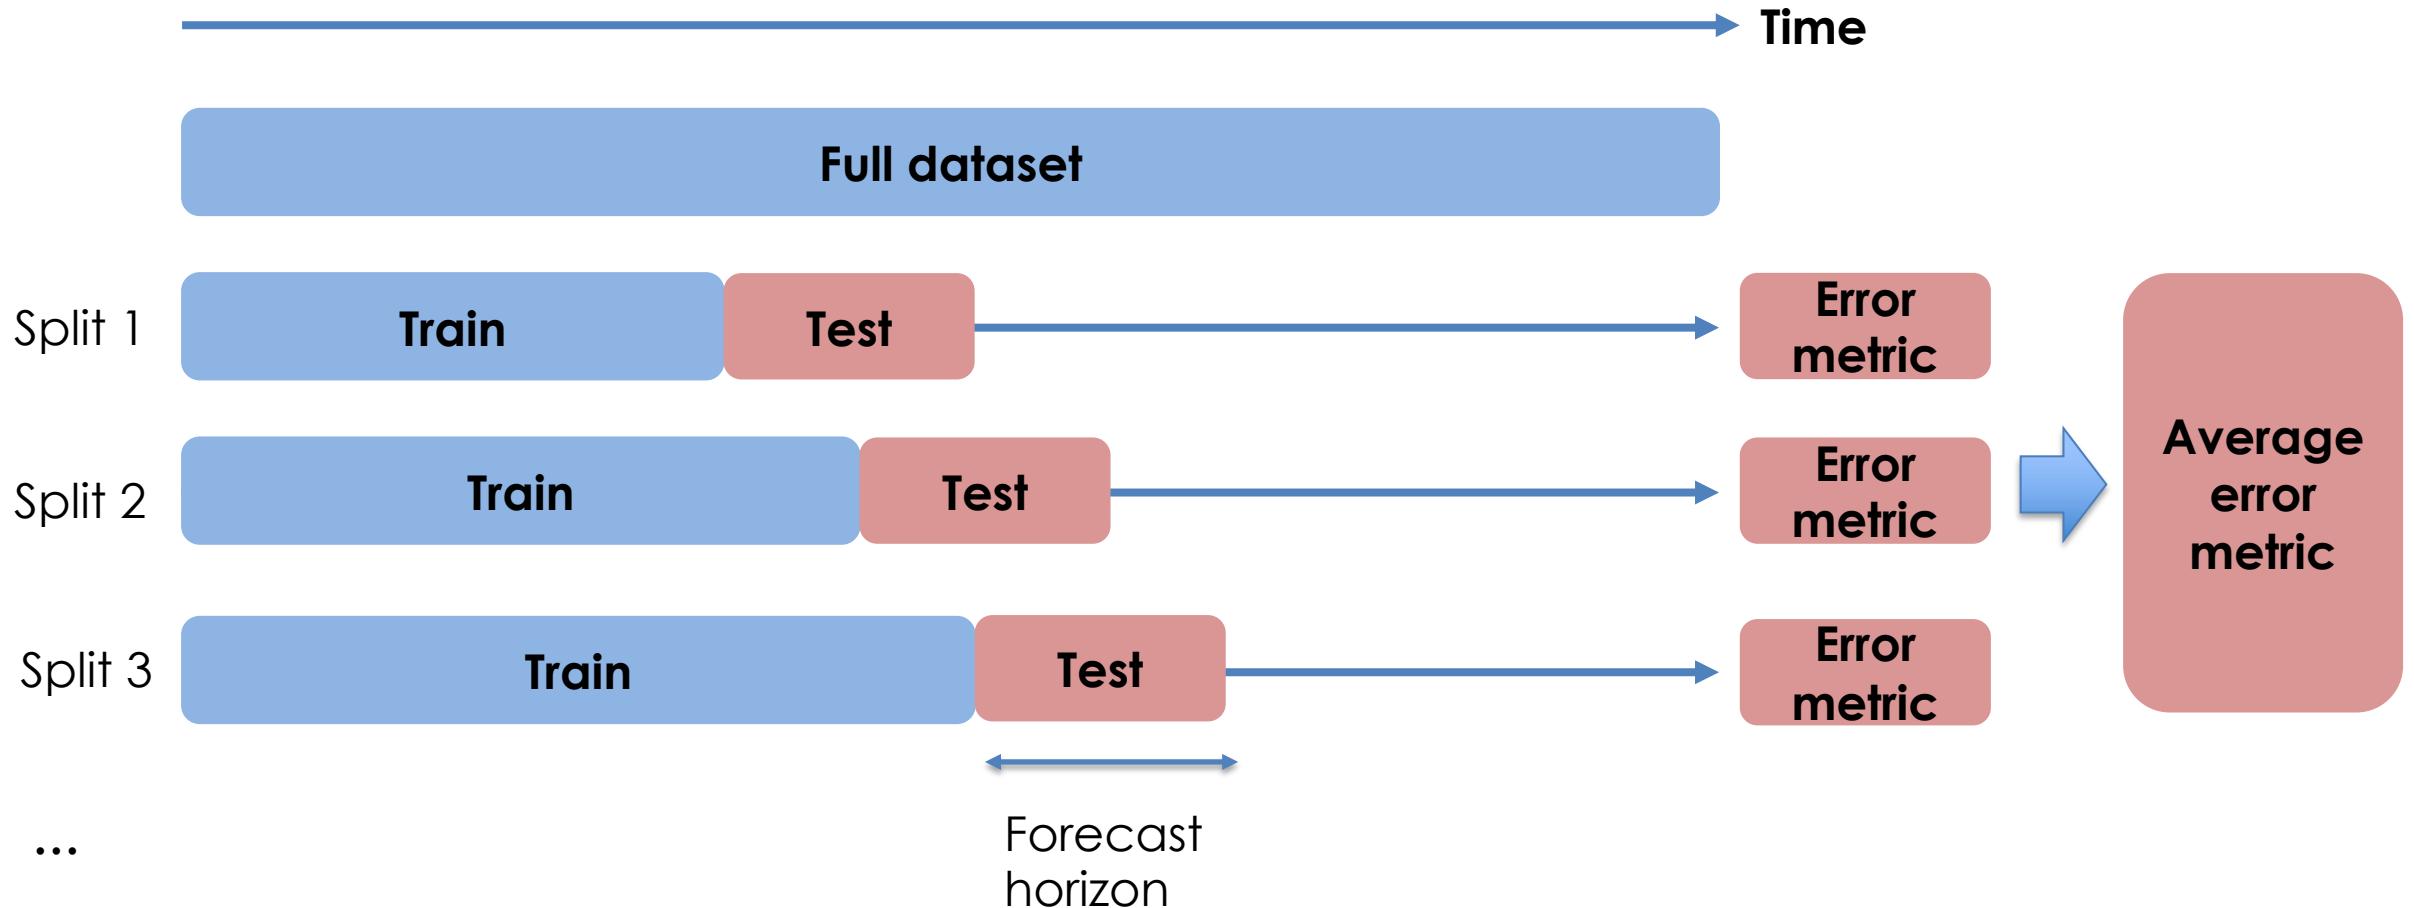
- Backtesting is where we **measure the performance** of a **forecasting model** on **historic data**.
- We want to **estimate the performance** of the model **on future data**.
- It is common to randomly shuffle data into train and test sets.
- We **can't randomly shuffle** the data into train and test sets when working **with time series**.

Train ■
Test ■

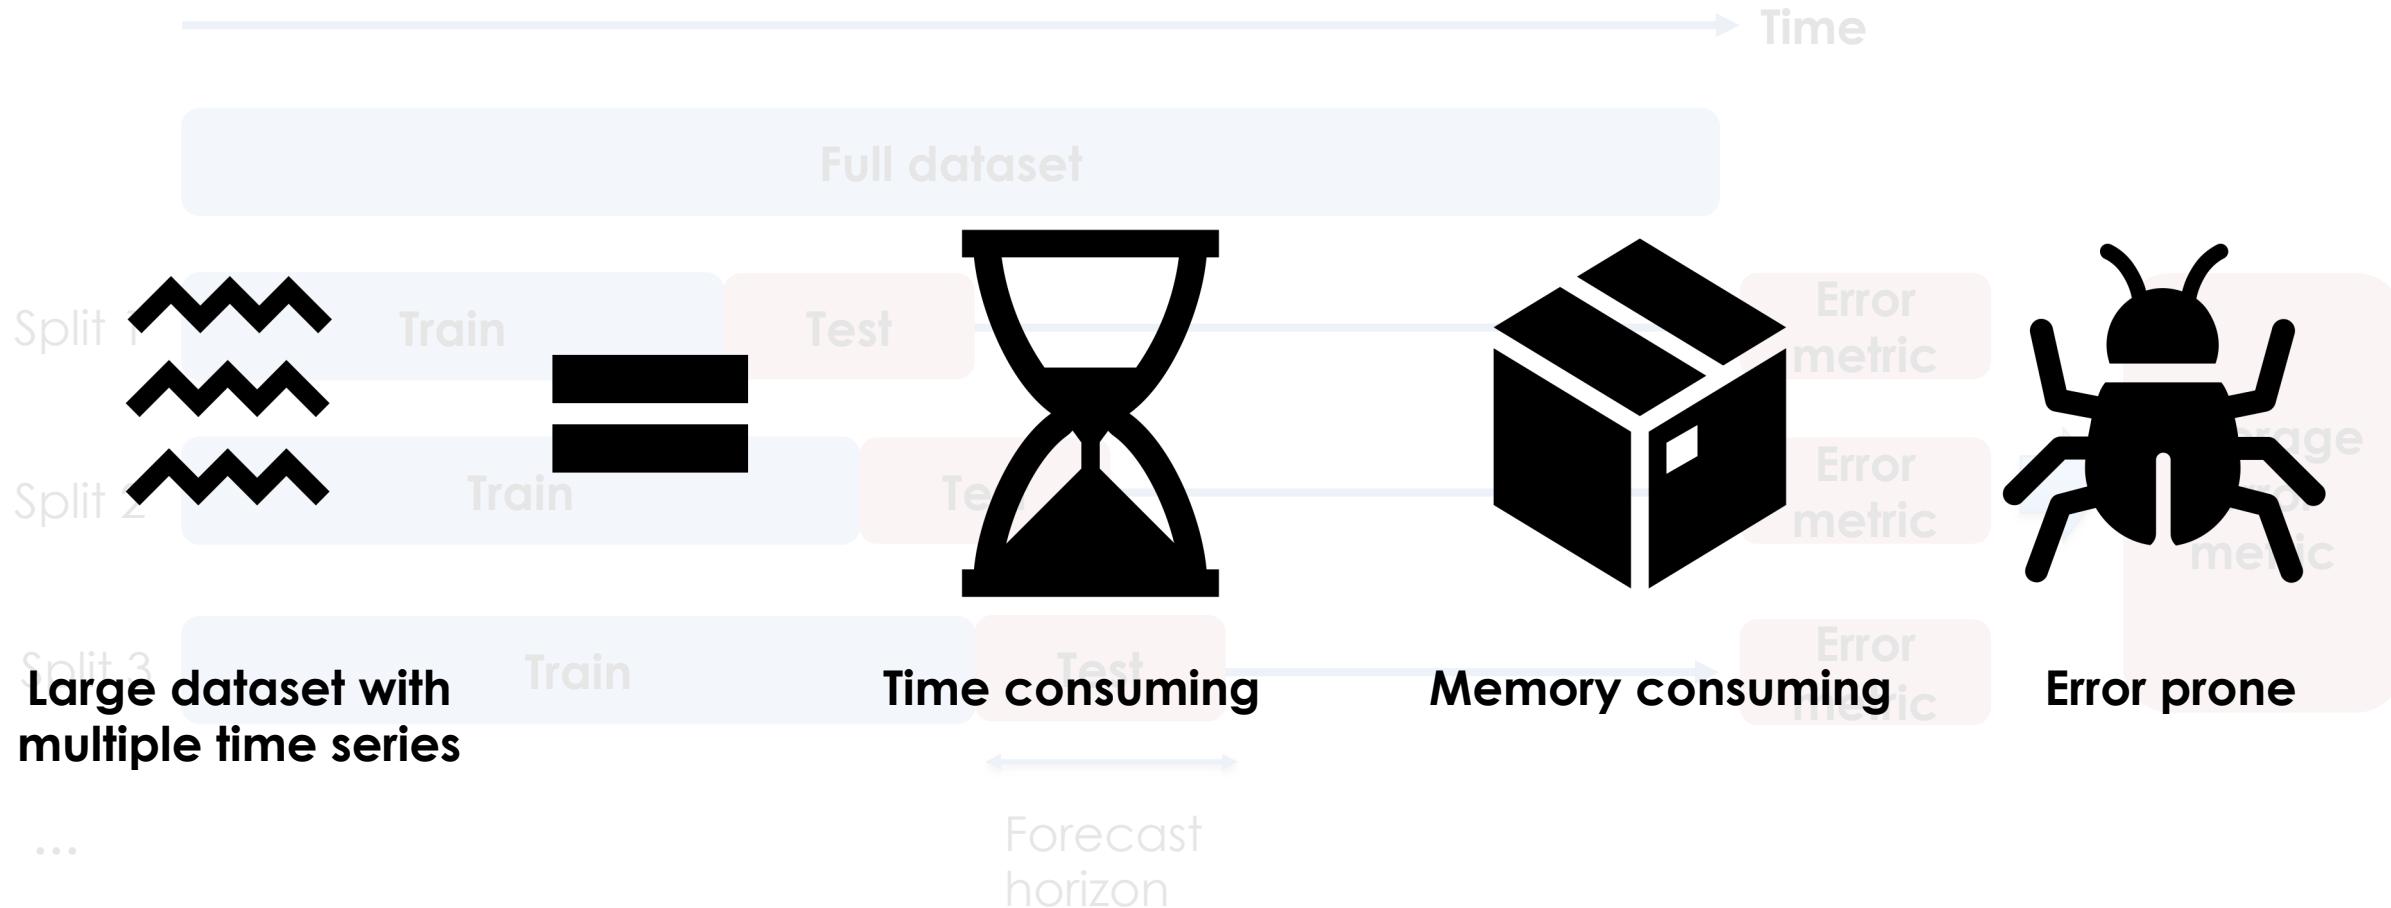


time	x1	x2	y
1	0
2	1
3	0
4	0
5	0
6	1
7	0

Backtesting



Backtesting



Backtesting parameters

Training
window
type

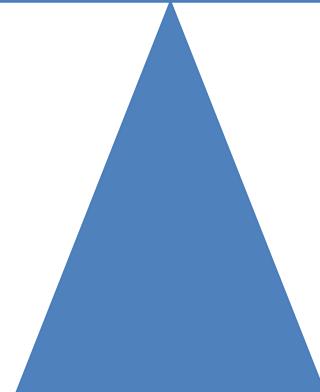
Step size

Refitting
frequency

Choosing backtesting parameters is a trade-off

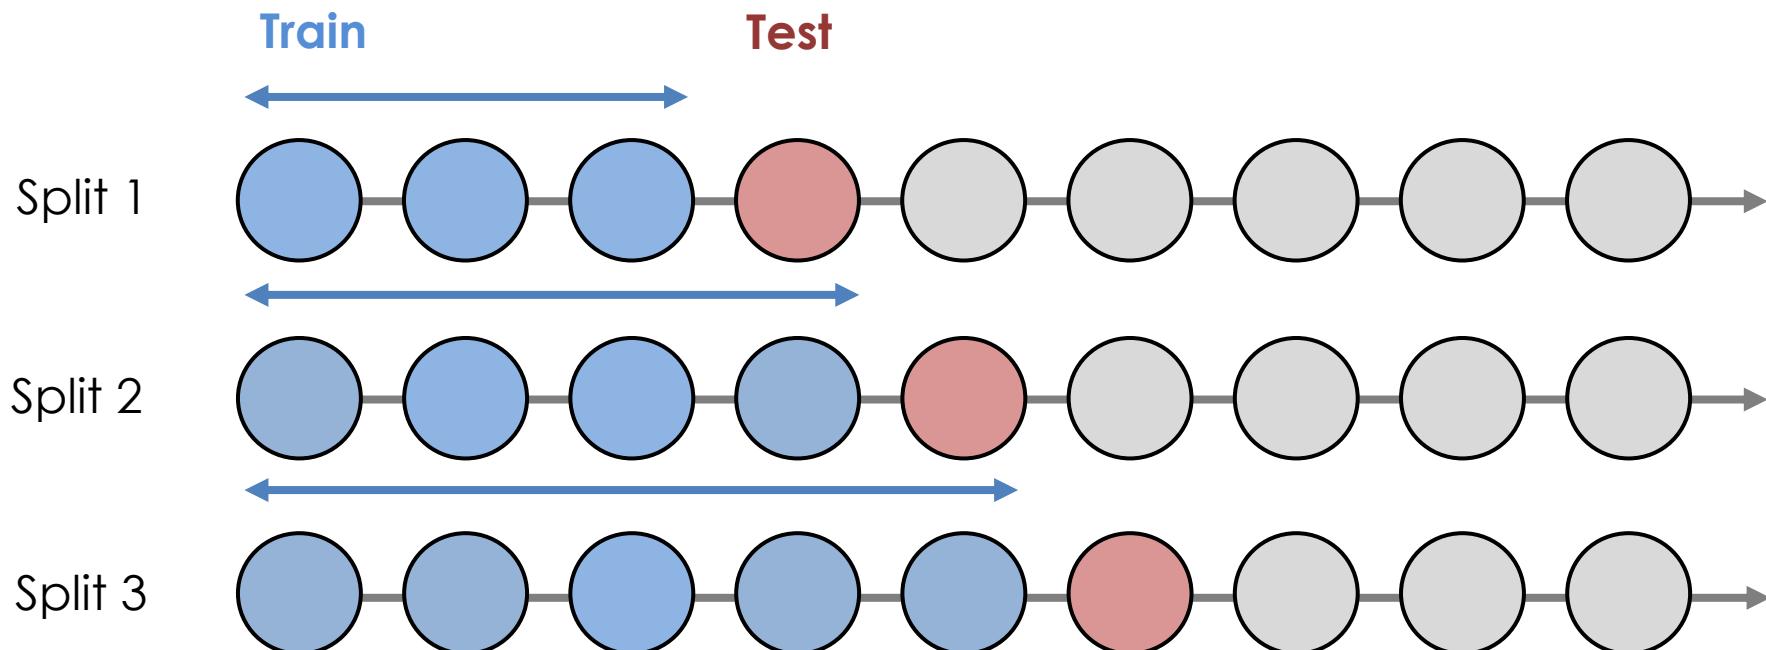
A reliable estimate
of model
performance.

Time and memory
consumption of
backtesting.



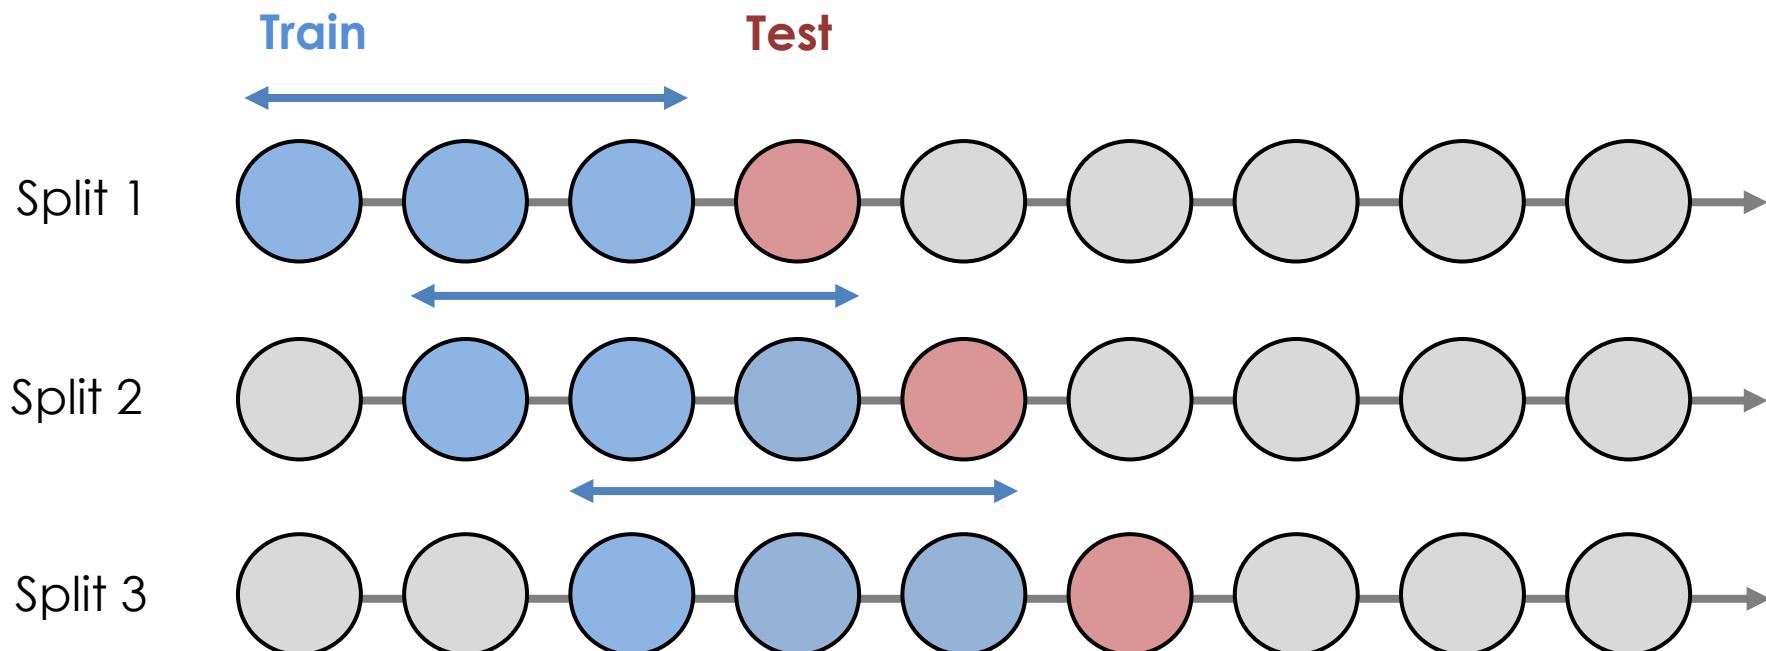
Backtesting parameters: Window type

- Expanding window



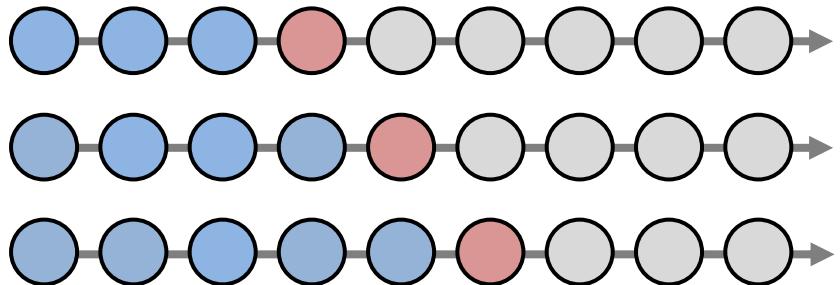
Backtesting parameters: Window type

- Rolling window (window size is fixed)



Backtesting parameters: Window type

Expanding window



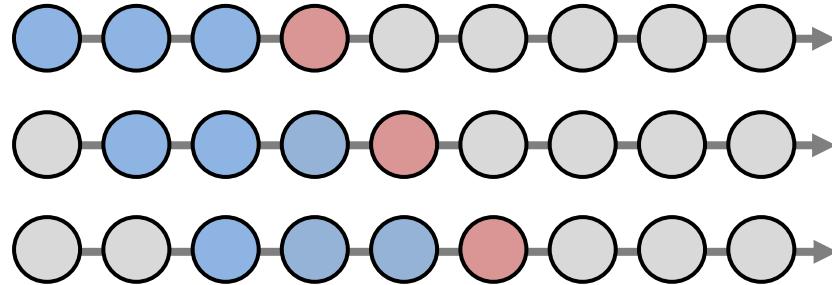
Pros

Uses all the training data.

Cons

More memory used.
Longer to train model.

Rolling window



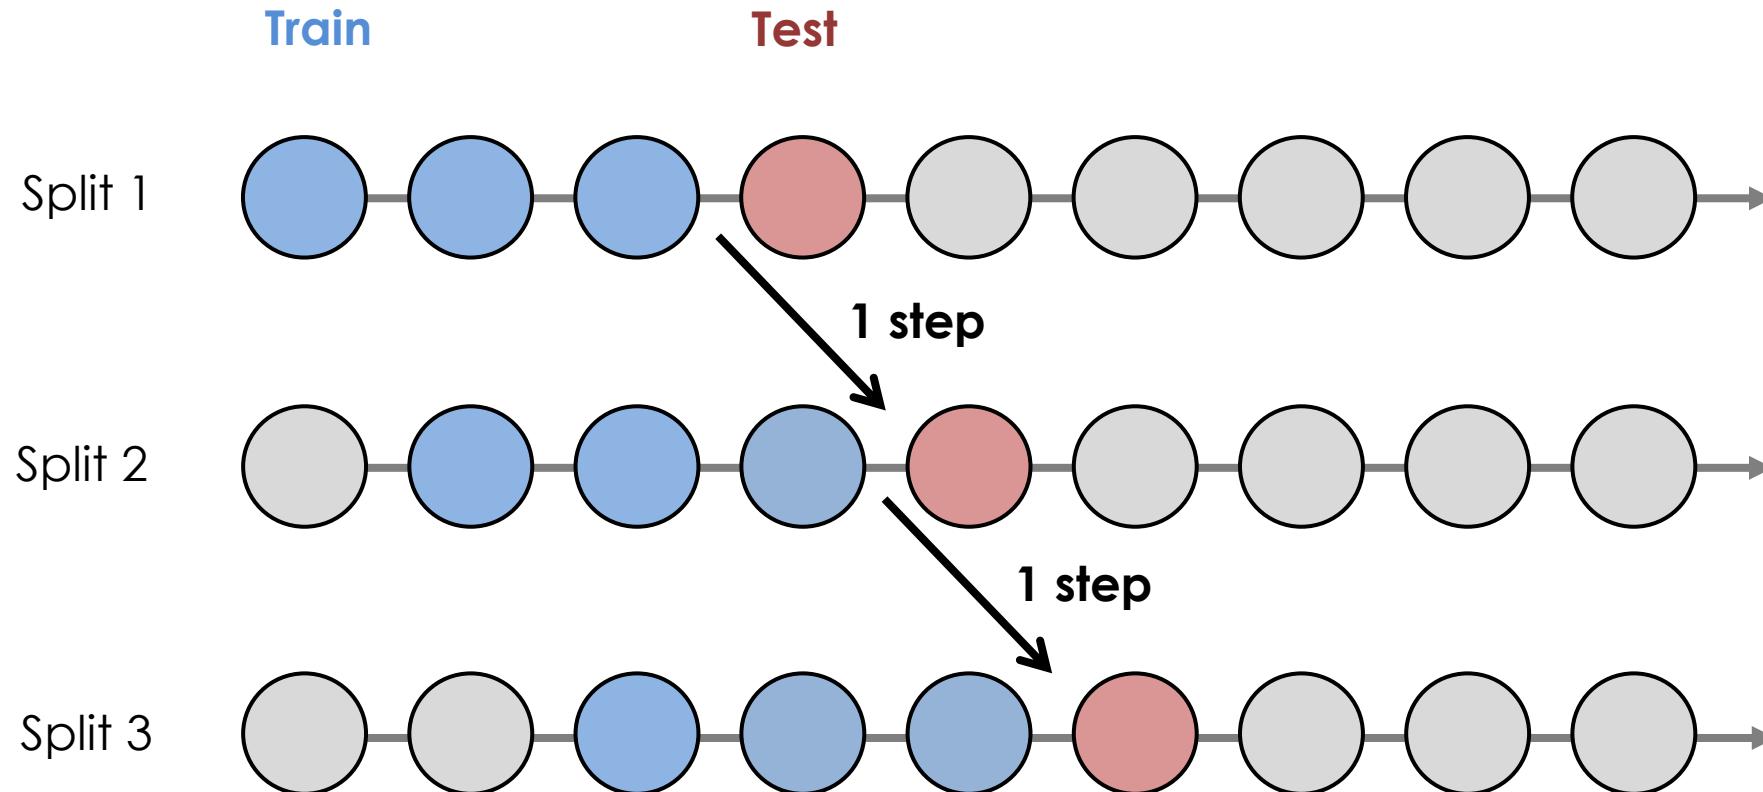
Pros

Less memory used.
Faster to train model.
Focuses more on recent data.

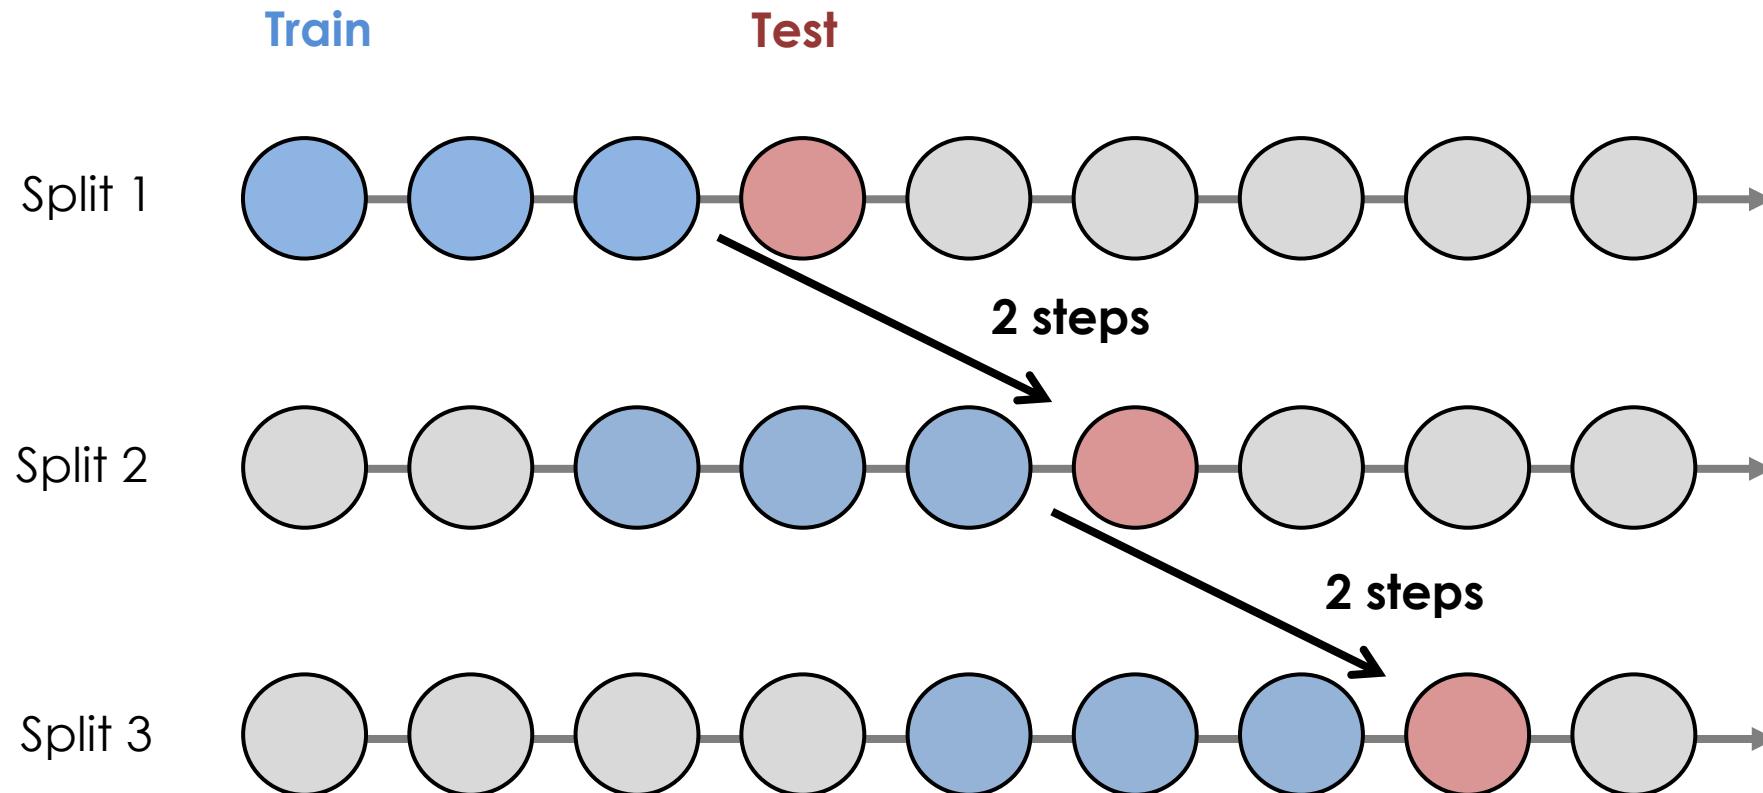
Cons

Less training data.

Backtesting parameters: Step size

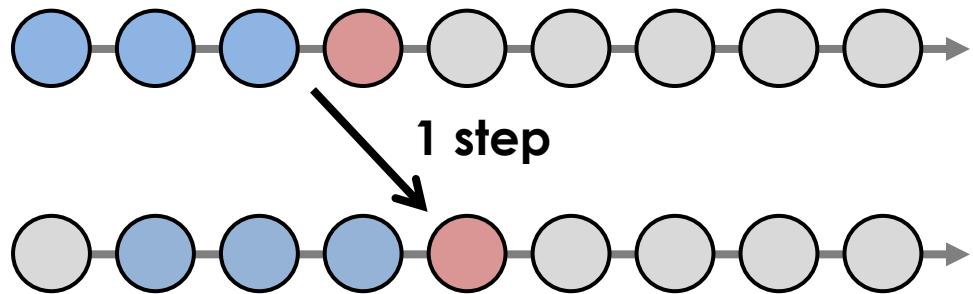


Backtesting parameters: Step size



Backtesting parameters: Step size

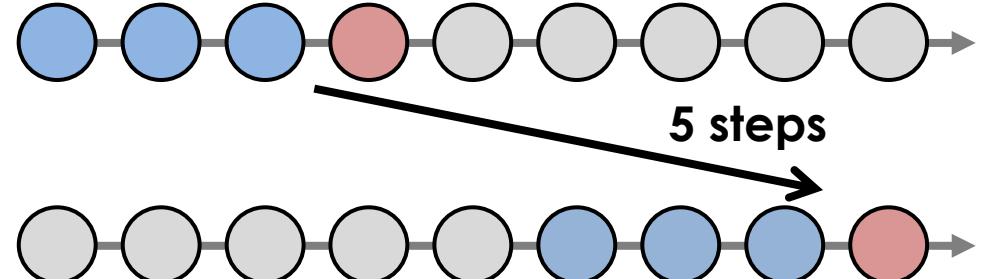
Small step size



Pros
Larger # of splits gives better error estimate.

Cons
More time and memory used.

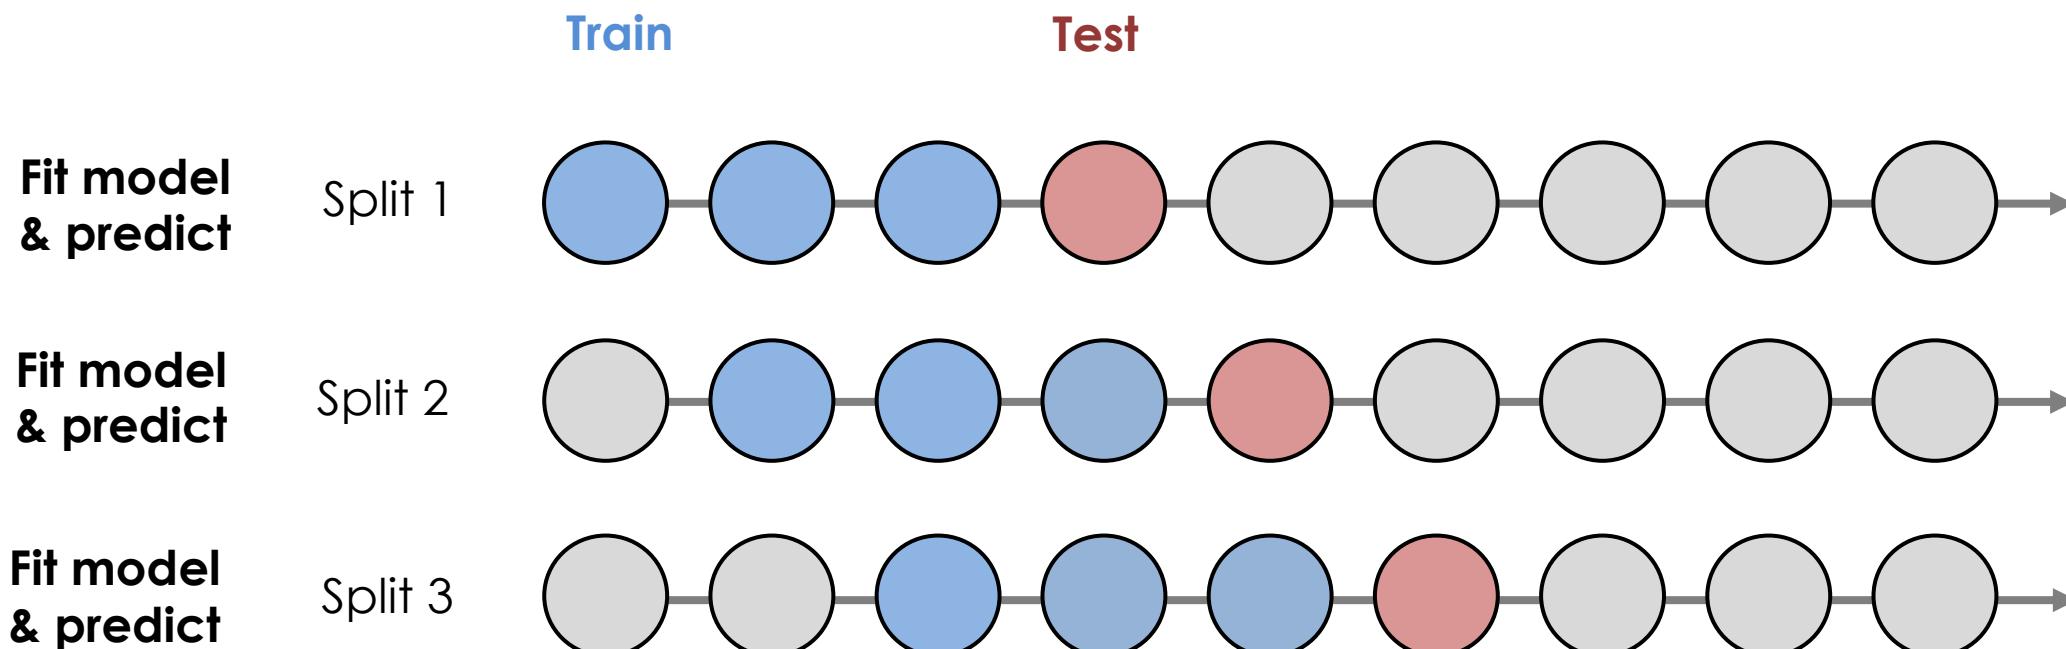
Large step size



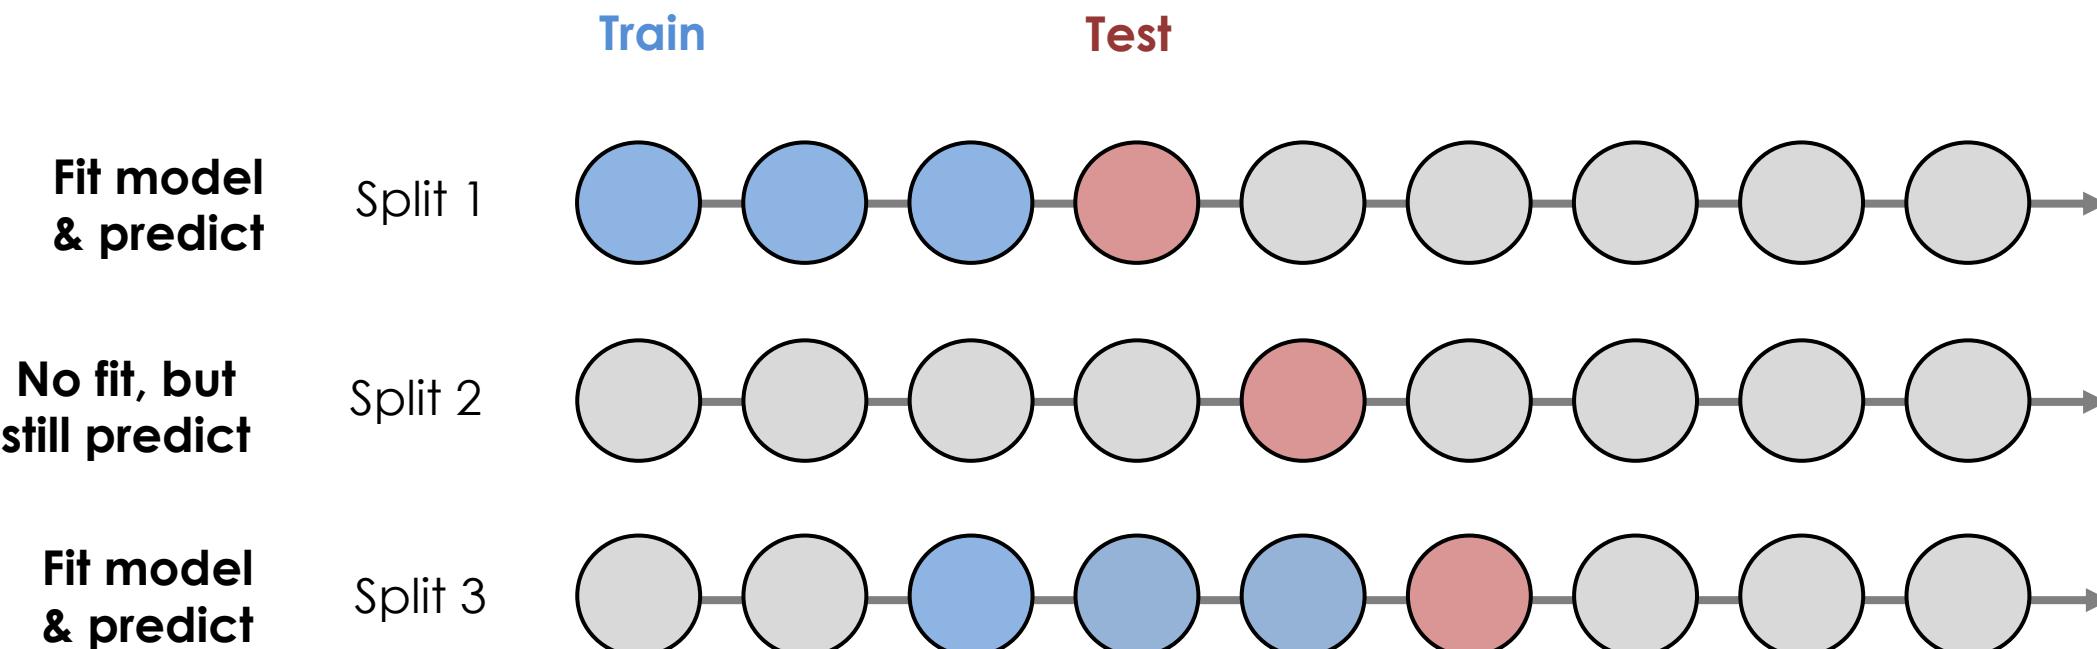
Pros
Less time and memory used.

Cons
Fewer # of splits gives worse error estimate.

Backtesting parameters: Refit frequency

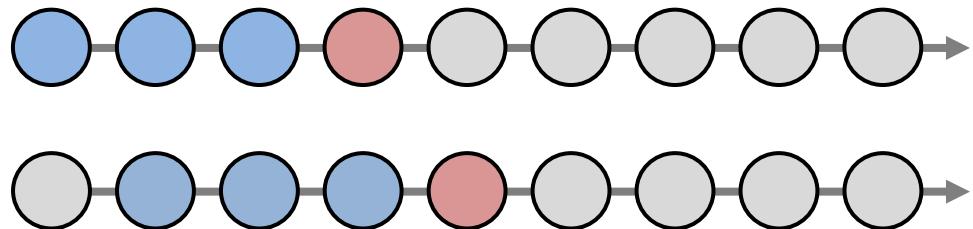


Backtesting parameters: Refit frequency

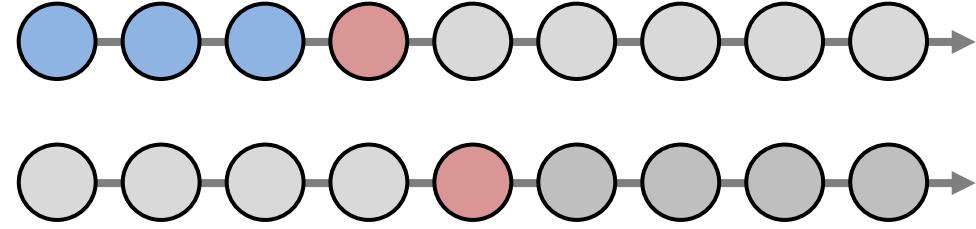


Backtesting parameters: Refit frequency

Refit every step



Refit periodically



Pros

Better estimate of model performance

Cons

More time

Pros

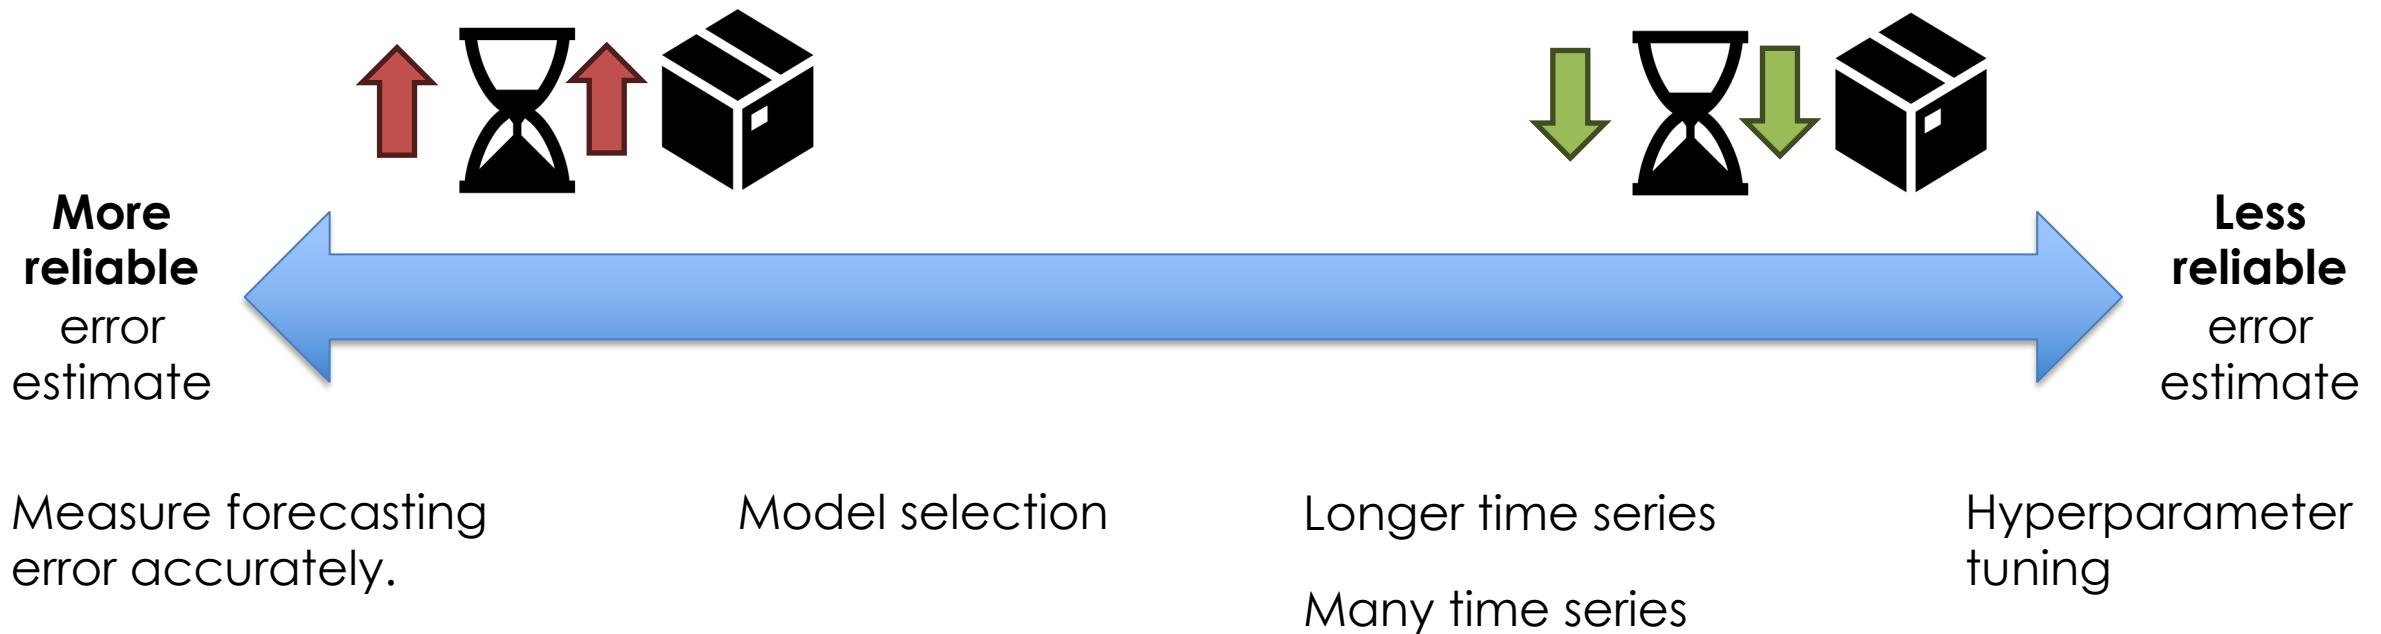
Less time

Cons

Worse estimate of model performance

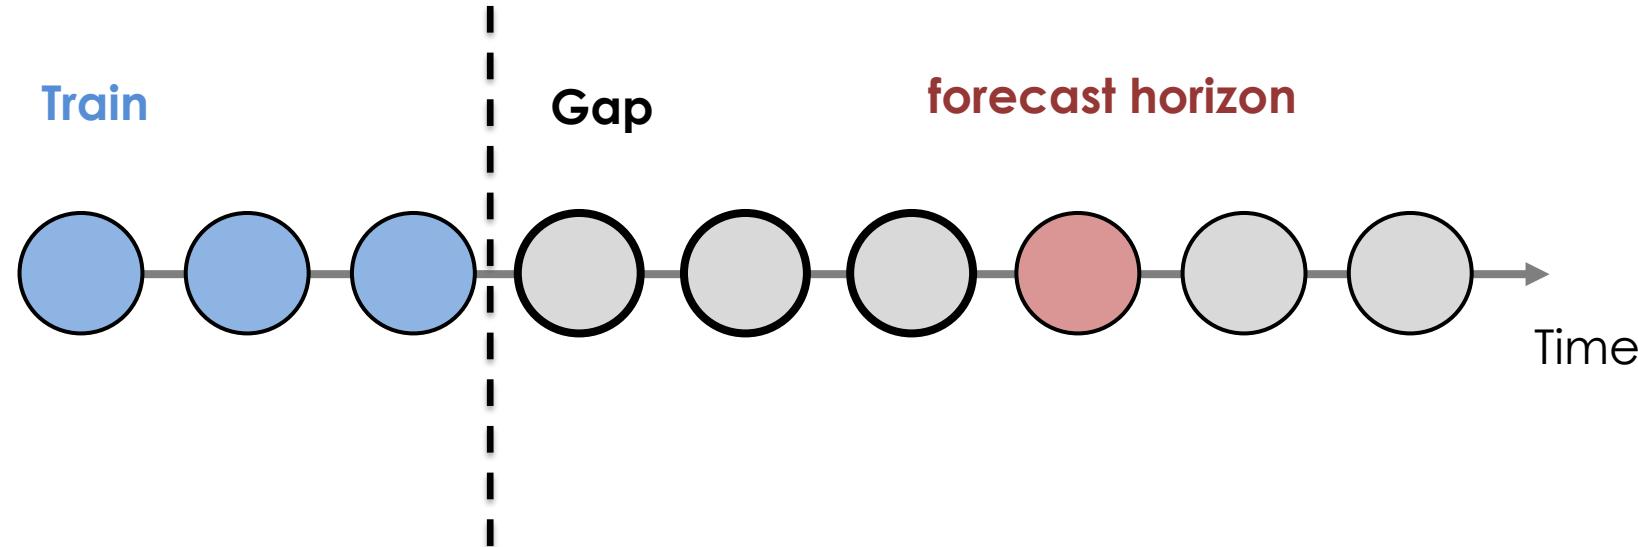
Choosing backtesting parameters is a trade-off

- We trade-off **reliability** against **time & memory**.
- The trade-off varies for **model selection**, **hyperparameter tuning**, and **error measurement**.



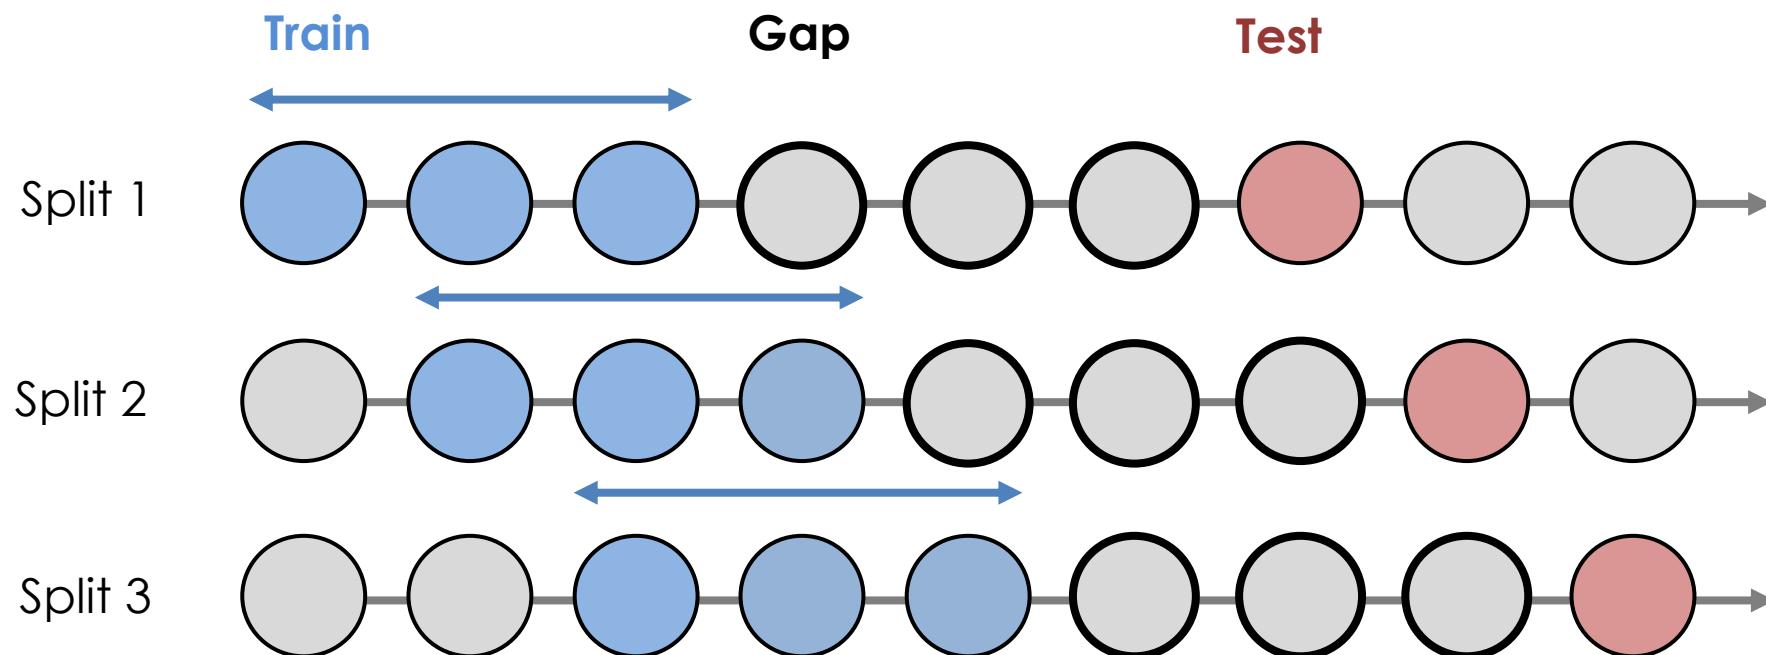
Backtesting: Try to reflect process in production

Example: There is a gap between when the forecast is created relative to when it is used. This gap should be reflected in backtesting.



Backtesting: Try to reflect process in production

Example: There is a gap between when the forecast is created relative to when it is used. This gap should be reflected in backtesting.



Backtesting: Edge cases for multiple series

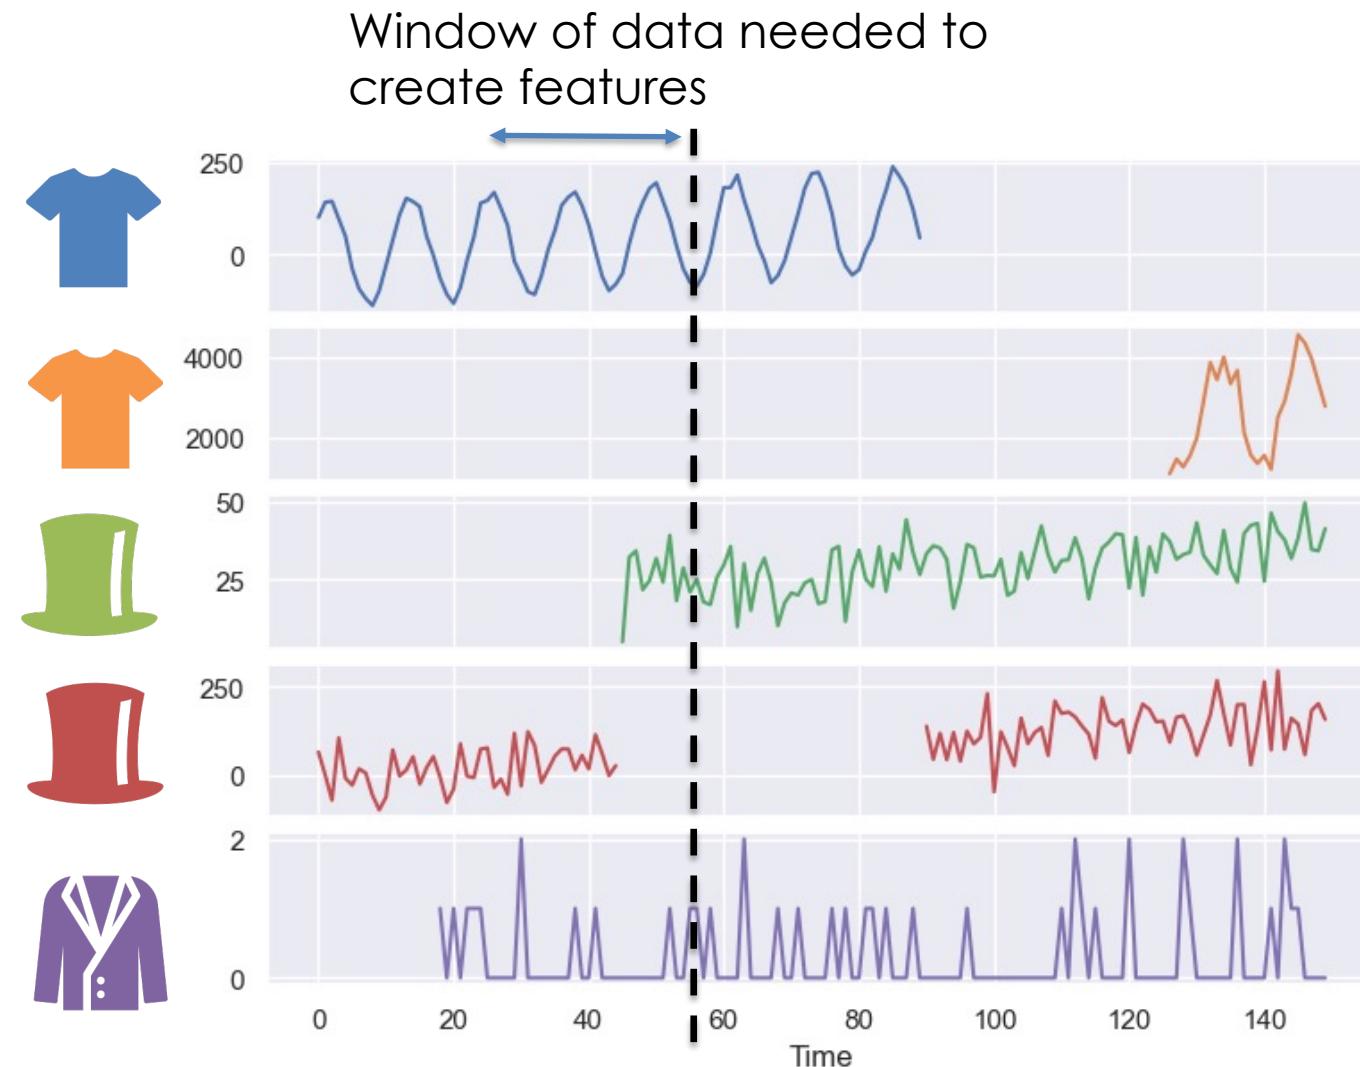
Product doesn't exist at train time but does exist at in forecast horizon.

Product does exist at train time but not in the the forecast horizon.



Backtesting: Edge cases for multiple series

Not enough historic data to create features.



Open source libraries make this easier

```
forecaster = ForecasterAutoreg(  
    regressor = RandomForestRegressor(random_state=123),  
    lags      = 15  
)  
  
metric, predictions = backtesting_forecaster(  
    forecaster          = forecaster,  
    y                   = data['y'],  
    steps               = 10,  
    metric              = 'mean_squared_error',  
    initial_train_size = 100,  
    fixed_train_size   = False,  
    gap                 = 0,  
    allow_incomplete_fold = True,  
    refit               = True,  
    n_jobs              = 'auto',  
    verbose             = True,  
    show_progress       = True  
)
```



Open source libraries make this easier

```
from functime.forecasting import linear_model
from functime.metrics import mase

forecaster = linear_model(lags=24, fit_intercept=False, freq="1mo")

y_preds, y_resids = forecaster.backtest(
    y=y_train,
    X=X_train,
    test_size=6,
    step_size=1,
    n_splits=3,
    window_size=1,
    strategy="expanding",
    drop_short=True,
)
```



Forecasting with machine learning in Python



 **functime**

 **NIXTLA**

Darts 

 **SKTIME**

Contents



Modern time series forecasting



Backtesting



Error metrics

There are a lot of forecasting error metrics

Table 8 Error measure definitions in the forecasting literature

Category	Error measure	Definition
Scale-Dependent Measures	Root Mean Squared Error (RMSE)	$RMSE = \sqrt{\frac{1}{n} \sum_{t=1}^n (e_t^2)}$
	Root Median Squared Error (RMdSE)	$RMdSE = \sqrt{\text{median}(e_t^2)}$
	Median Absolute Error (MdAE)	$MdAE = \text{median}(e_t)$
	Geometric Root Mean Squared Error (GRMSE, Syntetos and Boylan 2005)	$GRMSE = \sqrt[2n]{\prod_{t=1}^n e_t^2}$
	Geometric Mean Absolute Error (GMAE)	$GMAE = \sqrt[n]{\prod_{t=1}^n e_t }$
Measures based on Percentage Errors	Mean Absolute Percentage Error (MAPE)	$MAPE = \frac{1}{n} \sum_{t=1}^n (p_t)$

Hewamalage, Hansika, Klaus Ackermann, and Christoph Bergmeir. "Forecast evaluation for data scientists: common pitfalls and best practices." *Data Mining and Knowledge Discovery* 37.2 (2023): 788-832.

This review paper documents over 40 metrics.



Why are there so many error metrics?

How do we pick the right error metric(s)?

Structure of most error metrics

- The properties of an error metric depends on three main components that make a metric.

Mean absolute error

$$MAE = \text{mean}(|e_t|)$$

Structure of most error metrics

- The properties of an error metric depends on three main components that make a metric.
- Each forecast, \hat{y}_t , is associated with a **base error**.

Mean absolute error

$$MAE = \text{mean}(|e_t|)$$

↑
Scale-dependent base error:

$$e_t = y_t - \hat{y}_t$$

Structure of most error metrics

- The properties of an error metric depends on three main components that make a metric.
- Each forecast, \hat{y}_t , is associated with a **base error**.
- The base error is **transformed** to be **positive**.

Mean absolute error

$$MAE = \text{mean}(|e_t|)$$

Scale-dependent base error:

$$e_t = y_t - \hat{y}_t$$

Positive transform:

Absolute value: $|x|$

Structure of most error metrics

- The properties of an error metric depends on three main components that make a metric.
- Each forecast, \hat{y}_t , is associated with a **base error**.
- The base error is **transformed** to be **positive**.
- The base errors are **aggregated** with a **summary operator**.

Mean absolute error

$$MAE = \text{mean}(|e_t|)$$

Scale-dependent base error:

$$e_t = y_t - \hat{y}_t$$

Positive transform:

Absolute value: $|x|$

Summary operator:

mean

Structure of most error metrics

- The properties of an error metric depends on three main components that make a metric.
- Each forecast, \hat{y}_t , is associated with a **base error**.
- The base error is **transformed** to be **positive**.
- The base errors are **aggregated** with a **summary operator**.

Mean absolute percentage error

$$MAPE = \text{mean}(|p_t|)$$

Percentage error:

$$p_t = 100 \frac{y_t - \hat{y}_t}{y_t}$$

Positive transform:

Absolute value: $|x|$

Summary operator:

mean

Structure of most error metrics

- The properties of an error metric depends on three main components that make a metric.
- Each forecast, \hat{y}_t , is associated with a **base error**.
- The base error is **transformed** to be **positive**.
- The base errors are **aggregated** with a **summary operator**.
- Not all error metrics follow this pattern.

Mean absolute percentage error

$$MAPE = \text{mean}(|p_t|)$$

Percentage error:

$$p_t = 100 \frac{y_t - \hat{y}_t}{y_t}$$

Positive transform:

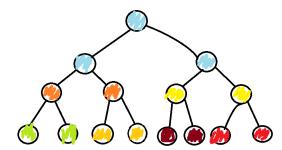
Absolute value: $|x|$

Summary operator:

mean

How to pick error metrics for forecasting?

Use case & properties of the time series



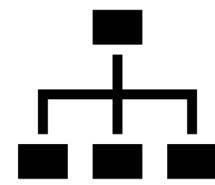
Modelling



Reporting



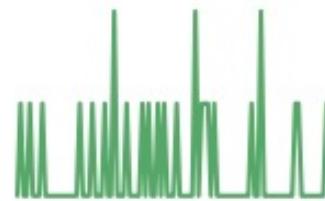
Does scale matter?



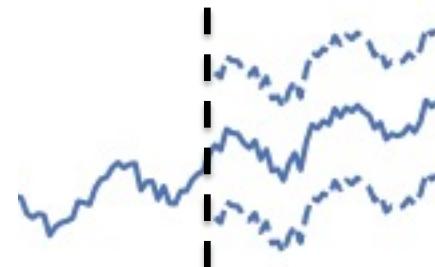
Hierarchy



Level shifts



Intermittency



Over vs under forecasting

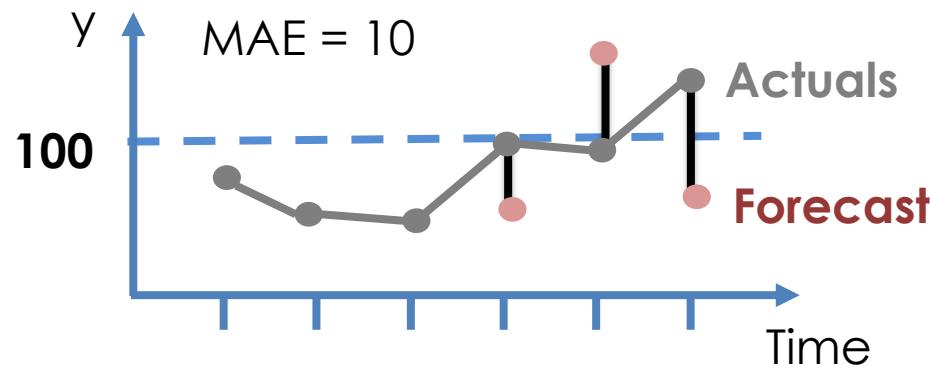


Outliers

Properties of error metrics

$$E(y_t, \hat{y}_t)$$

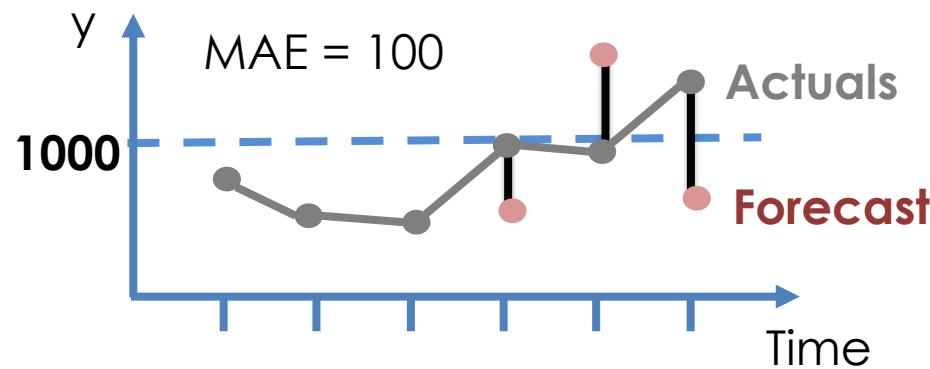
How to pick error metrics for forecasting?



Properties of error metrics

$$E(y_t, \hat{y}_t)$$

- **Scale dependence.**

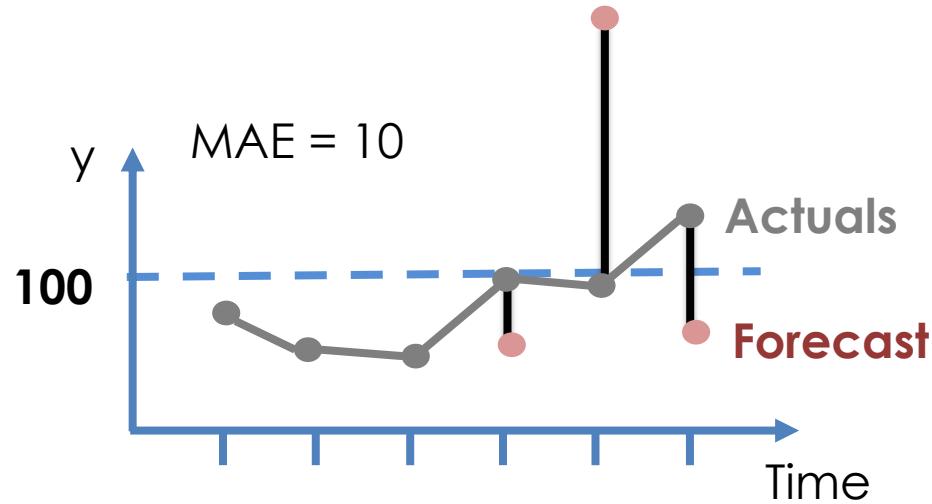


How to pick error metrics for forecasting?

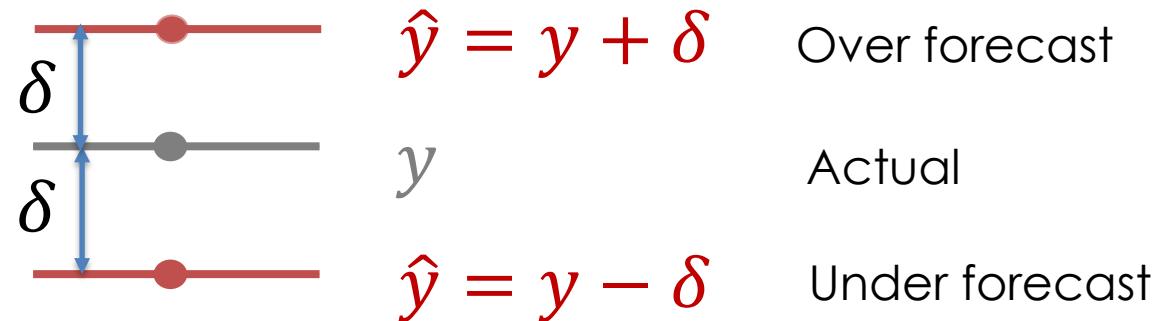
Properties of error metrics

$$E(y_t, \hat{y}_t)$$

- Scale dependence.
- **Sensitive to outliers.**



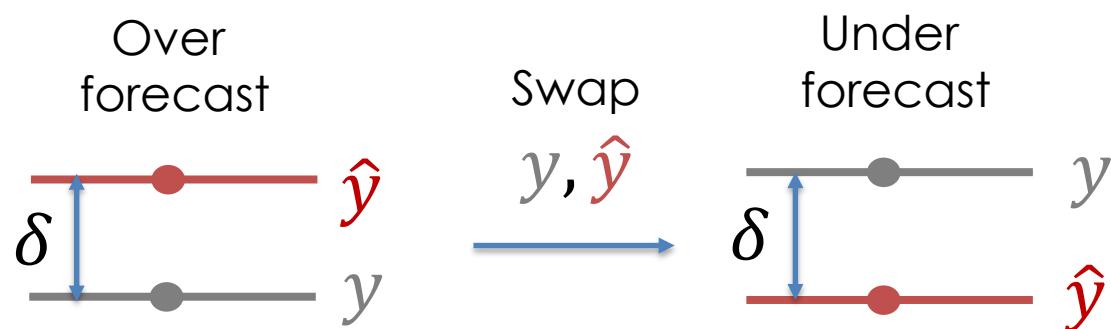
How to pick error metrics for forecasting?



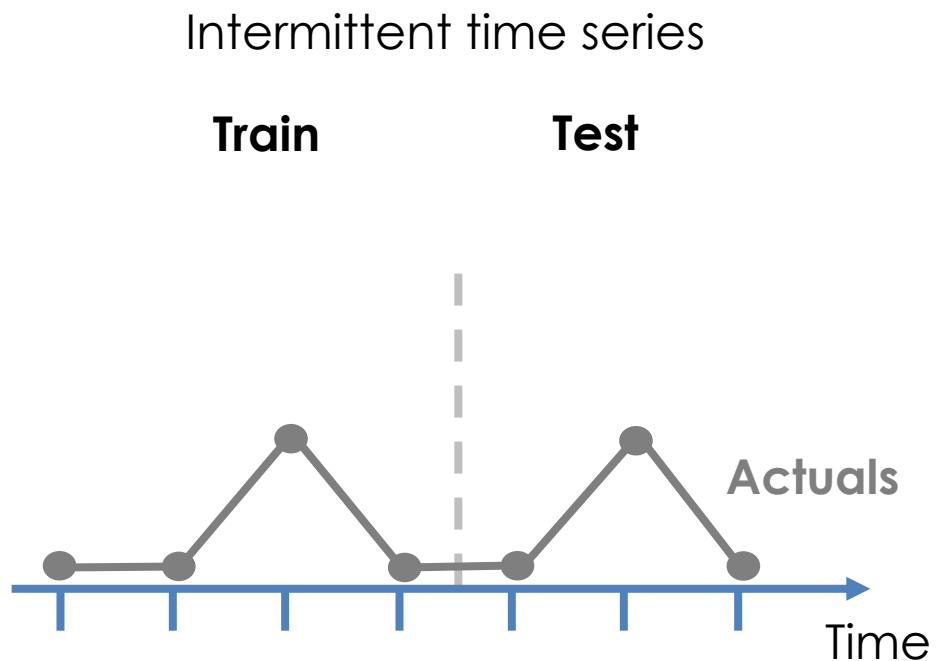
Properties of error metrics

$$E(y_t, \hat{y}_t)$$

- Scale dependence.
- Sensitive to outliers.
- **Symmetric to over and under forecasting.**



How to pick error metrics for forecasting?



Properties of error metrics

$$E(y_t, \hat{y}_t)$$

- Scale dependence.
- Sensitive to outliers.
- Symmetric to over and under forecasting.
- **Can it handle zeros?**

How to pick error metrics for forecasting?

Properties of error metrics

The **RMSE**:

$$\sqrt{\frac{1}{N} \sum_t (y_t - \hat{y}_t)^2}$$

is **minimized by the mean**:

$$\hat{y}_t = \text{mean}(y_t)$$

The **MAE**:

$$\frac{1}{N} \sum_t |y_t - \hat{y}_t|$$

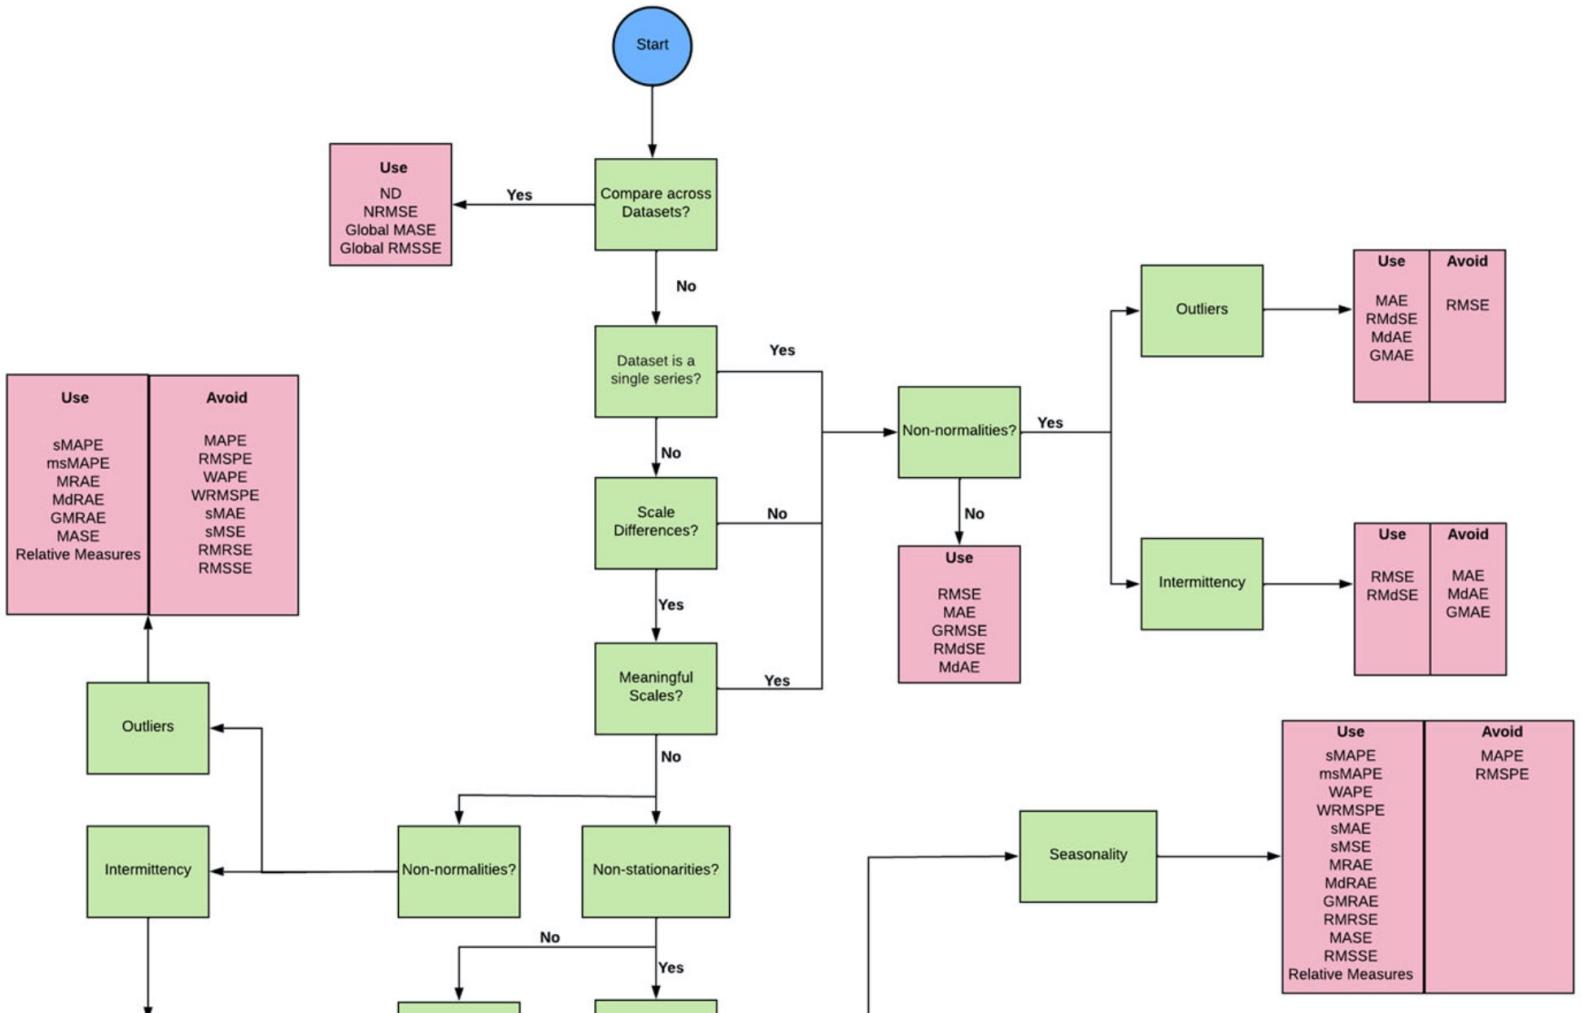
is **minimized by the median**:

$$\hat{y}_t = \text{median}(y_t)$$

$$E(y_t, \hat{y}_t)$$

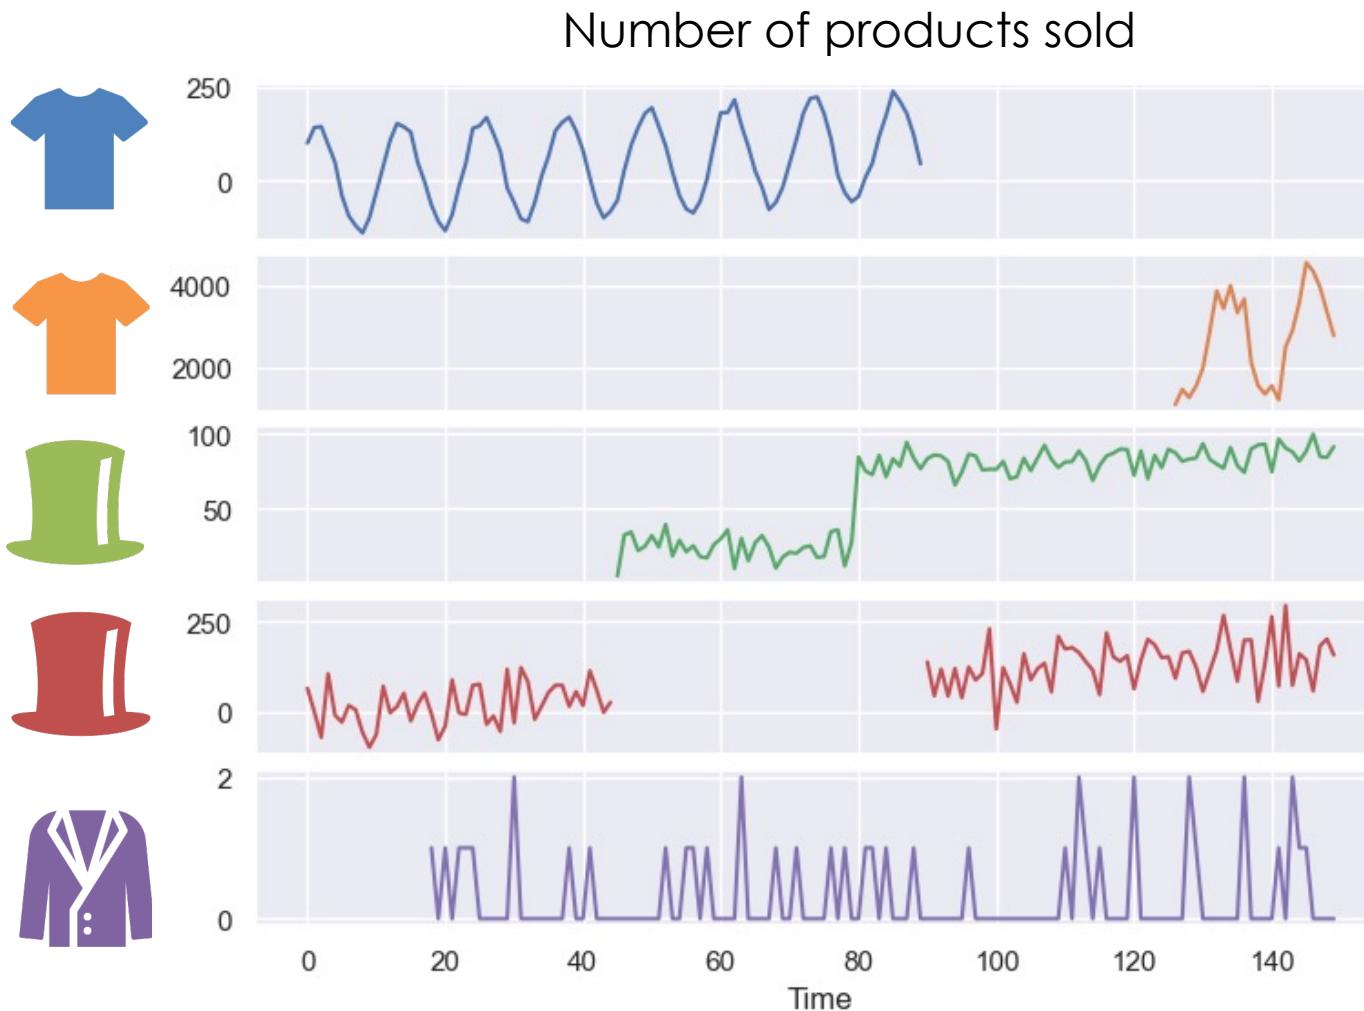
- Scale dependence.
- Sensitive to outliers.
- Symmetric to over and under forecasting.
- Can it handle zeros?
- **Optimise for mean, median, etc.**

Guidelines



Hewamalage, Hansika, Klaus Ackermann, and Christoph Bergmeir. "Forecast evaluation for data scientists: common pitfalls and best practices." *Data Mining and Knowledge Discovery* 37.2 (2023): 788-832.

Multiple time series where the scale matters



We want:

- Larger time series to contribute more to the error (i.e., scale dependence).
- Symmetry to over/under forecasting.
- Tolerate zeros.
- To be able to compare errors for different subsets of timeseries at different scales.
- RMSE and MAE cover **almost** all these requirements.

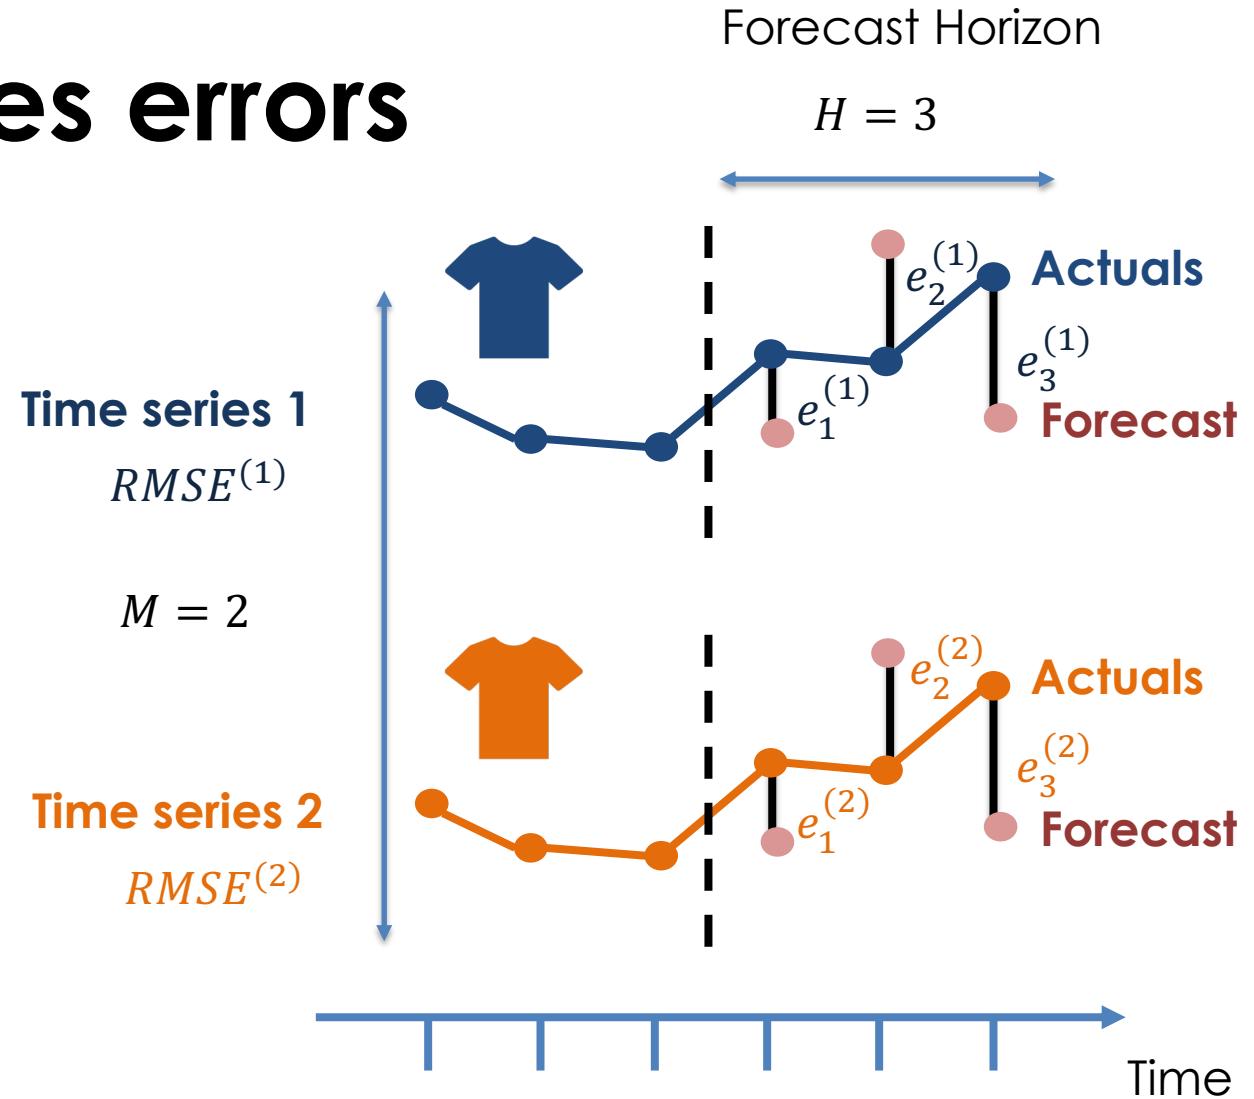
Multiple time series errors

Averaging

$$\text{mean RMSE} = \frac{1}{M} \sum_{i=1}^M \text{RMSE}^{(i)}$$

Pooling

$$\text{RMSE} = \sqrt{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H (e_t^{(i)})^2}$$



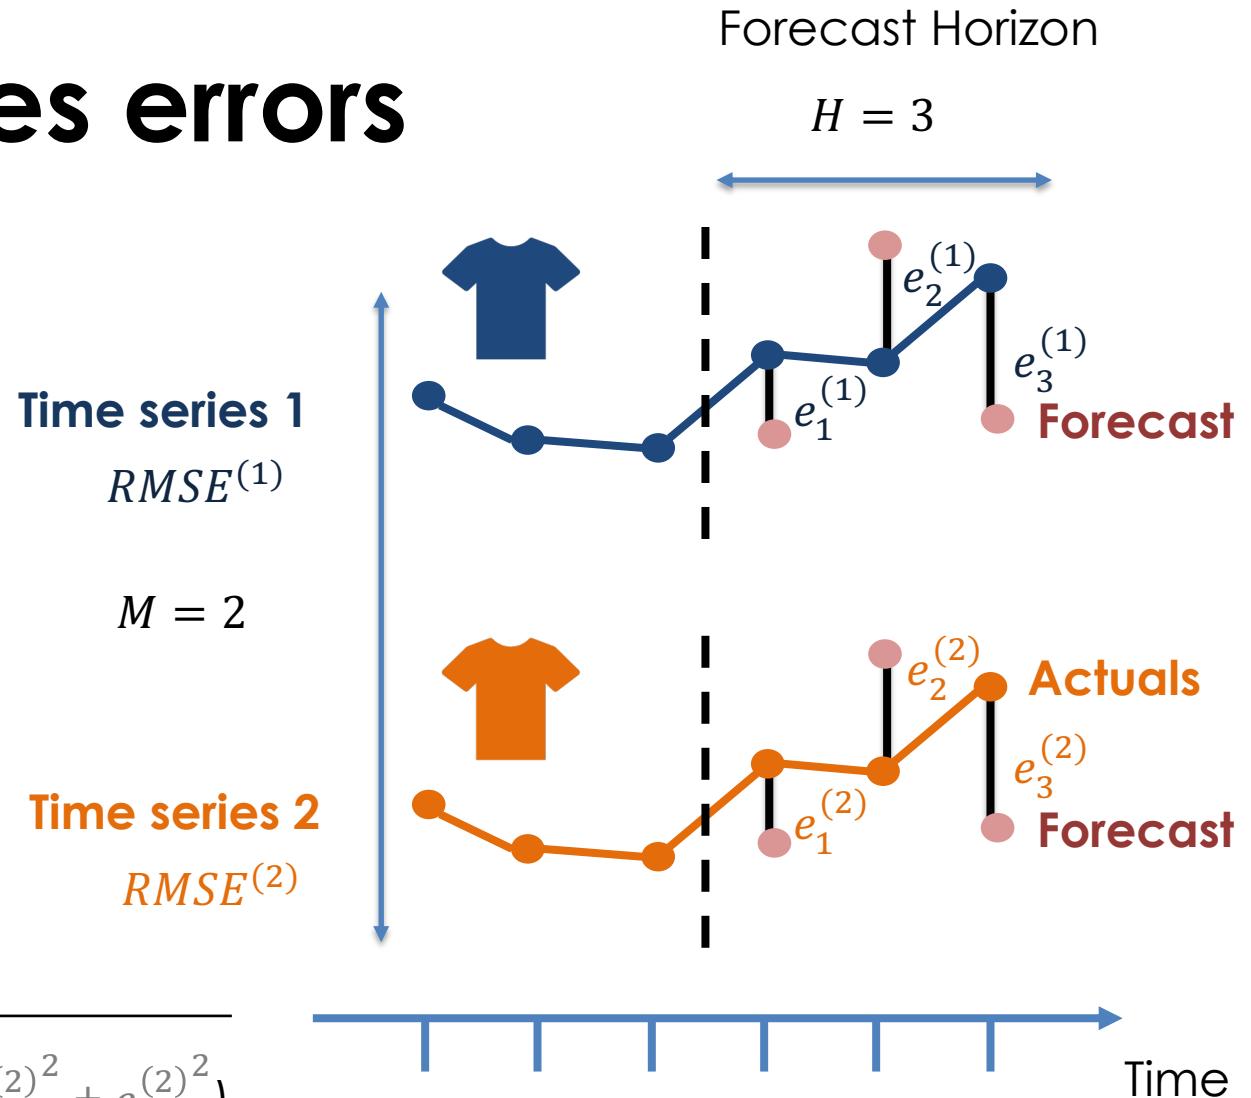
Multiple time series errors

Averaging

$$\text{mean RMSE} = \frac{1}{2}(\text{RMSE}^{(1)} + \text{RMSE}^{(2)})$$

Pooling

$$\text{RMSE} = \sqrt{\frac{1}{2 * 3}(e_1^{(1)2} + e_2^{(1)2} + e_3^{(1)2} + e_1^{(2)2} + e_2^{(2)2} + e_3^{(2)2})}$$



Multiple time series errors

Pool **scale-dependent errors** over multiple time series and **scale by mean of all the time series**.

Normalised RMSE

$$NRMSE = \frac{\sqrt{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H (e_t^{(i)})^2}}{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |y_t^{(i)}|} = \frac{RMSE}{\bar{|y|}}$$

Normalised Deviation
(aka normalized MAE)

$$ND = \frac{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |e_t^{(i)}|}{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |y_t^{(i)}|} = \frac{MAE}{\bar{|y|}}$$

Multiple time series errors

Larger time series contribute more to the error than smaller ones.

Normalised RMSE

$$NRMSE = \frac{\sqrt{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H (e_t^{(i)})^2}}{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |y_t^{(i)}|} = \frac{RMSE}{\bar{|y|}}$$

Normalised Deviation
(aka normalized MAE)

$$ND = \frac{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |e_t^{(i)}|}{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |y_t^{(i)}|} = \frac{MAE}{\bar{|y|}}$$

Multiple time series errors

Symmetric to over/under forecasting and can tolerate zeros.

Normalised RMSE

$$NRMSE = \frac{\sqrt{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H (e_t^{(i)})^2}}{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |y_t^{(i)}|} = \frac{RMSE}{\bar{|y|}}$$

Normalised Deviation
(aka normalized MAE)

$$ND = \frac{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |e_t^{(i)}|}{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |y_t^{(i)}|} = \frac{MAE}{\bar{|y|}}$$

Multiple time series errors

These metrics can be used to **compare errors between different subsets/datasets of time series**.

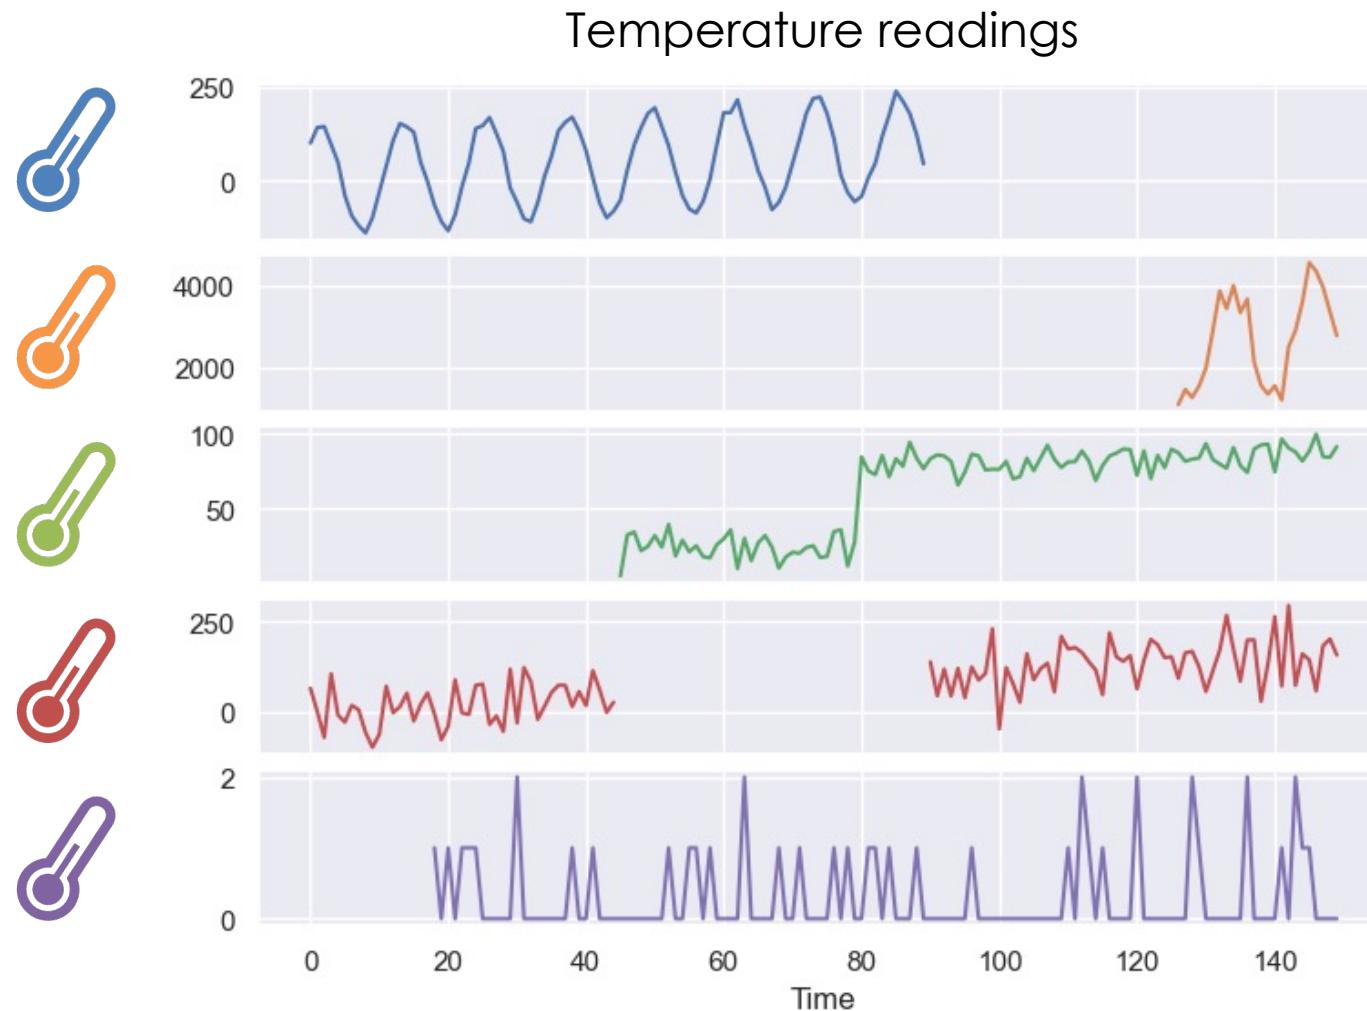
Normalised RMSE

$$NRMSE = \frac{\sqrt{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H (e_t^{(i)})^2}}{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |y_t^{(i)}|} = \frac{RMSE}{\bar{|y|}}$$

Normalised Deviation
(aka normalized MAE)

$$ND = \frac{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |e_t^{(i)}|}{\frac{1}{MH} \sum_{i=1}^M \sum_{t=1}^H |y_t^{(i)}|} = \frac{MAE}{\bar{|y|}}$$

Multiple time series where the scale is arbitrary

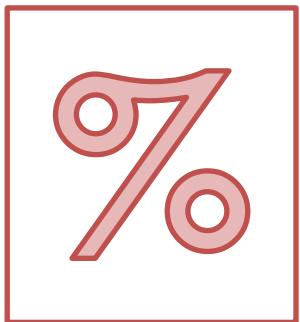


We want:

- Each time series to contribute equally to the error regardless of scale.
- Symmetry to over/under forecasting.
- Tolerate zeros.
- To be able to compare errors for different subsets of timeseries.
- **We need to use scale-independent error metrics.**

Scale-independent error metrics

Percentage error



$$p_t = \frac{e_t}{y_t}$$

Weighted mean absolute percentage error

$$WAPE = \frac{1}{\sum_t w_t} \sum_{t=1}^H w_t |p_t|$$

$$\xrightarrow{w_t=|y_t|} = \frac{1}{\sum_t |y_t|} \sum_{t=1}^H |e_t| \\ = \frac{MAE}{mean(|y_t|)}$$

Scale independent when comparing series. ✓

Interpretable. ✓

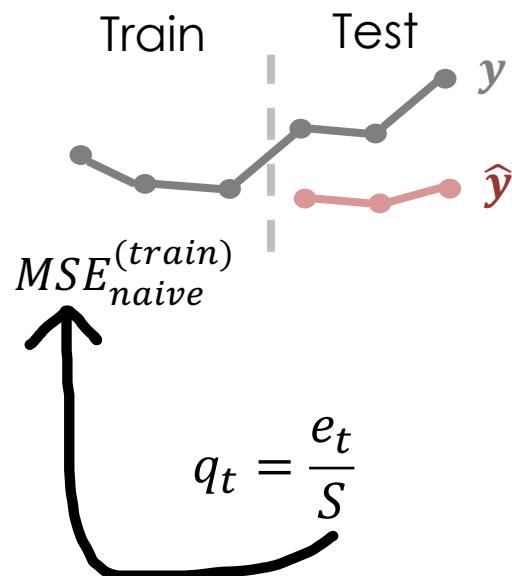
Symmetric to over/under forecasting ✓

Well defined when $\hat{y}_t = 0$ or some $y_t = 0$. ✓

Not well defined when all $y_t = 0$ in horizon. ✗

Scale-independent error metrics

Scaled error



Mean absolute scaled error

$$MASE = \text{mean}(|q_t|)$$

$$S = MAE_{naive}^{(train)}$$

Root mean squared scaled error

$$RMSSE = \sqrt{\text{mean}(q_t^2)}$$

$$S = MSE_{naive}^{(train)}$$

Scale independent. ✓

Symmetric to over/under forecasting ✓

Well defined when $\hat{y}_t = 0$ or $y_t = 0$. ✓

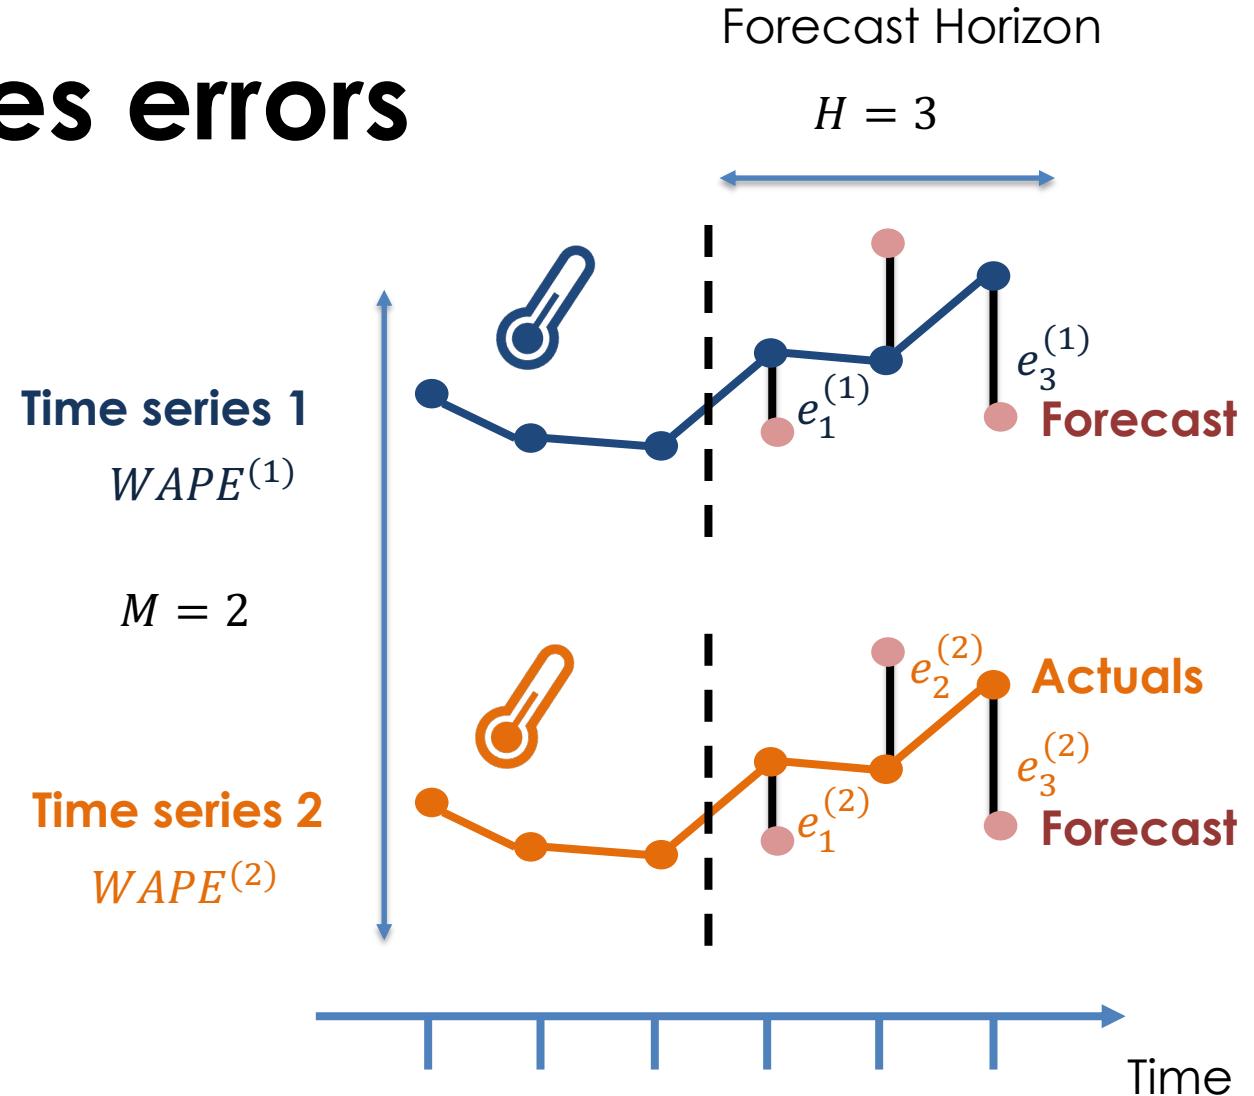
Not well defined when all $y_t = 0$ in train set. ✗

Multiple time series errors

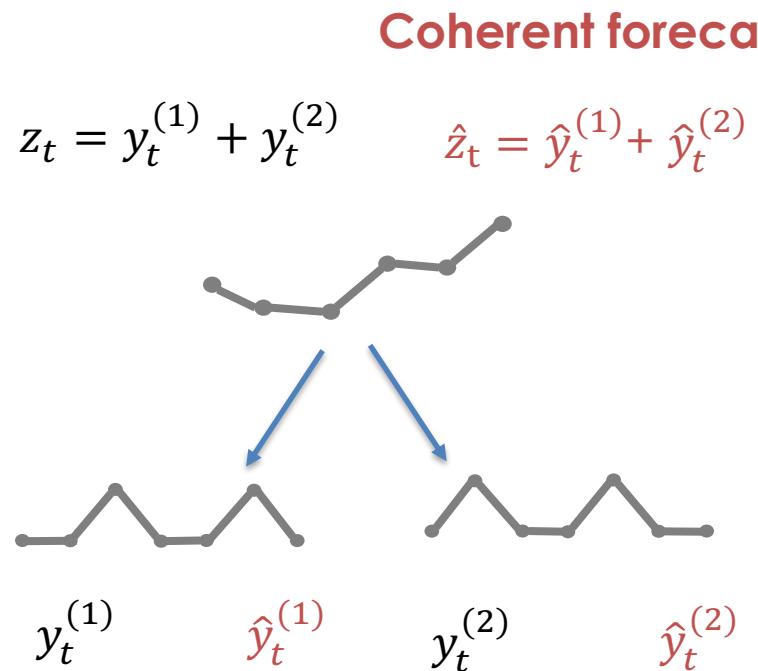
Average the scale-independent error metrics

$$\text{mean WAPE} = \frac{\sum_{i=1}^M \text{WAPE}^{(i)}}{M}$$

$$= \frac{1}{2} (\text{WAPE}^{(1)} + \text{WAPE}^{(2)})$$



Hierarchical forecasts



We want the most accurate forecasts at all levels.

We pick an error metric E .

We create a forecast to minimise $E(y_t^{(1)}, \hat{y}_t^{(1)})$ and $E(y_t^{(2)}, \hat{y}_t^{(2)})$.

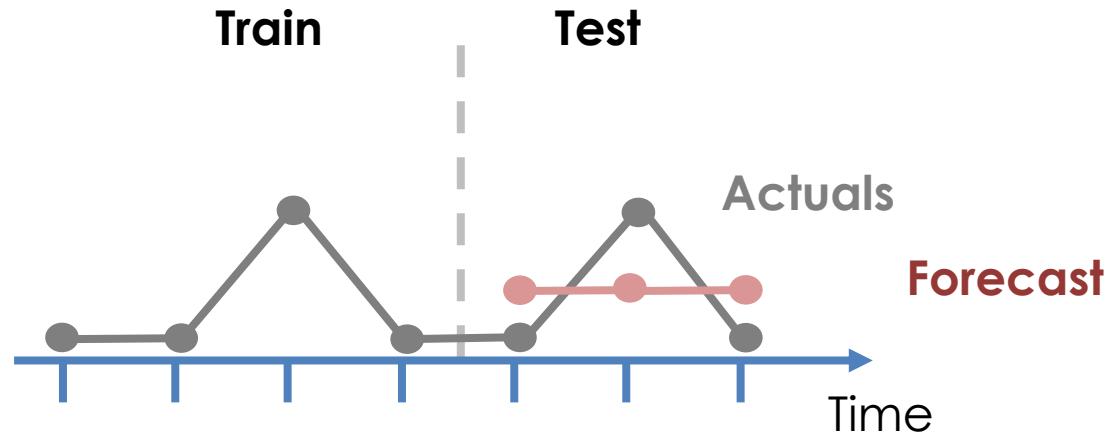
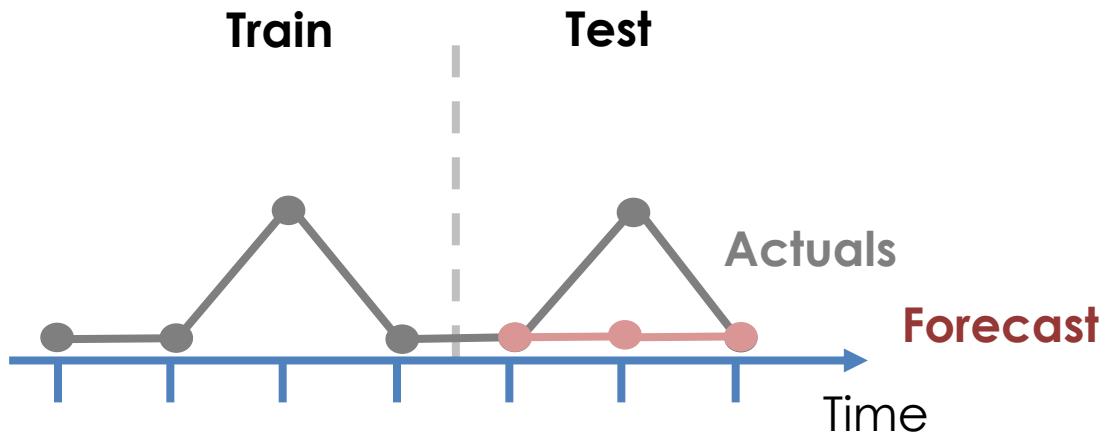
Does this also minimize $E(z_t, \hat{z}_t)$?

No! Only if the error metric is a **function of e_t^2** .

Examples: RMSE, RMSSE

Kolassa, S., 2023. Do we want coherent hierarchical forecasts, or minimal MAPEs or MAEs?(We won't get both!). *International Journal of Forecasting*, 39(4), pp.1512-1517.

Intermittent time series



- MAE is minimised by the median.
- “Best” forecast is zero. MAE is bad metric for intermittent data.

- RMSE is minimised by the mean.
- “Best” forecast is a non-zero constant: “avg. units sold over time”.

Conclusions

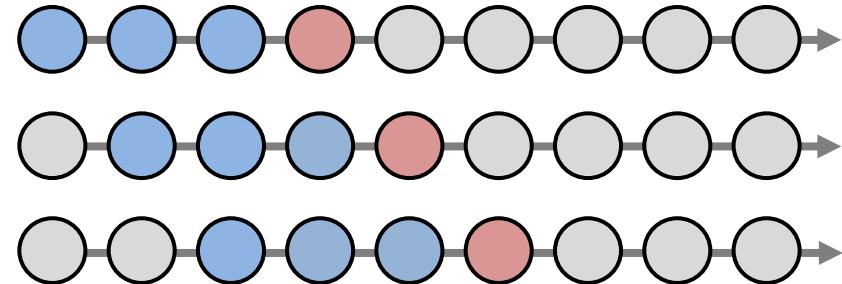
1. **Modern time series forecasting** involves many related time series.



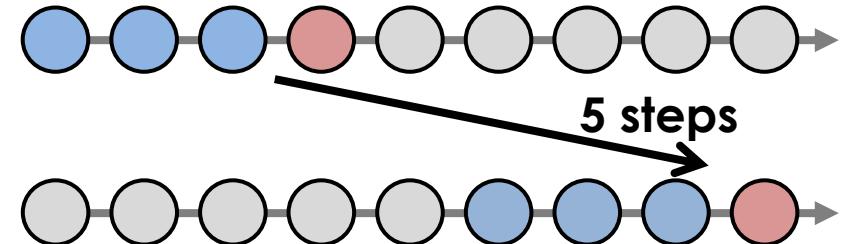
Conclusions

1. **Modern time series forecasting** involves many related time series.
2. Selecting **backtesting** parameters is a trade-off between reliability and speed.

Rolling window

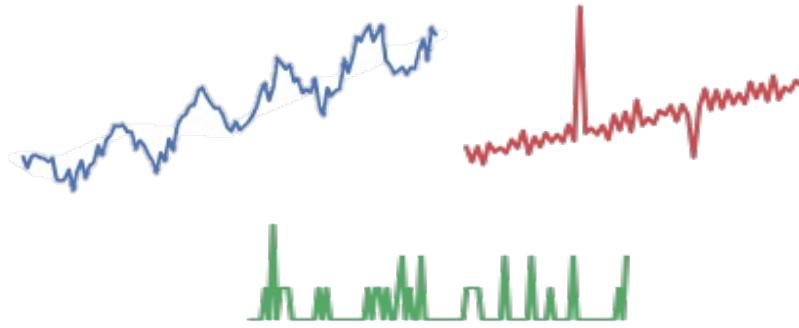


Step size

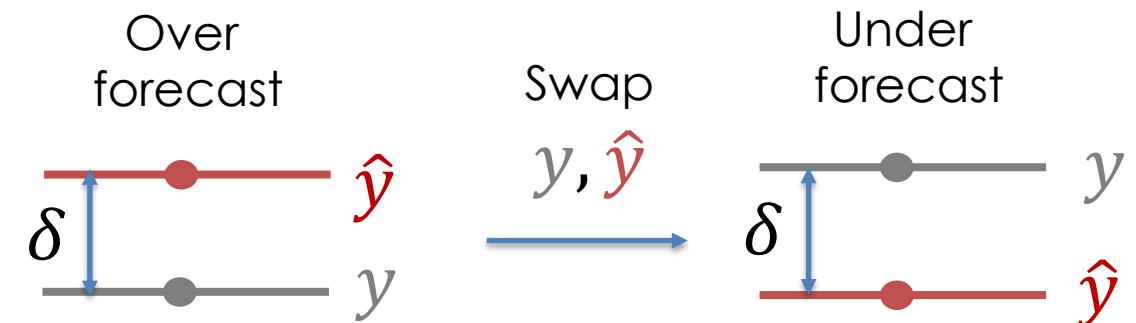


Conclusions

1. **Modern time series forecasting** involves many related time series.
2. Selecting **backtesting** parameters is a trade-off between reliability and speed.
3. Selecting **error metrics** requires pairing properties of our data and error metrics.



$$MAE = \frac{1}{N} \sum_i |e_i|$$



Conclusions

1. Modern time series forecasting involves many related time series.
2. Selecting backtesting parameters is a trade-off between reliability and speed.
3. Selecting error metrics requires pairing properties of our data and error metrics.
4. Many python libraries are available to support with forecasting.



If you'd like to learn more ...

trainindata.com/p/forecasting-specialization



In/kishanmanani



@KishManani



In/soledad-galli

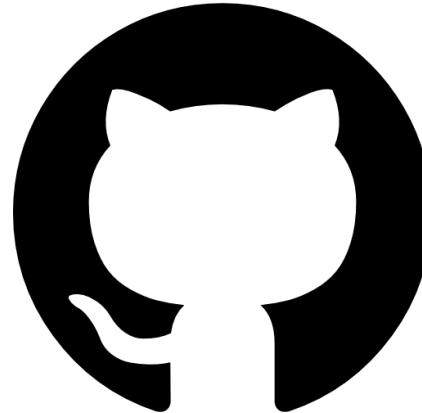


@Soledad_Galli

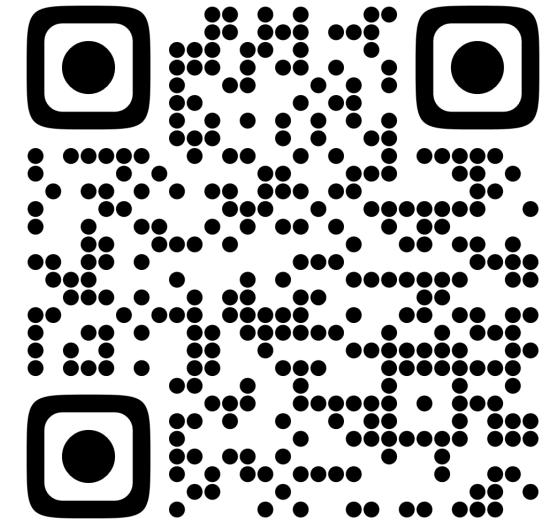
References

- [1] Hewamalage, H., Ackermann, K. and Bergmeir, C., 2023. Forecast evaluation for data scientists: common pitfalls and best practices. *Data Mining and Knowledge Discovery*, 37(2), pp.788-832.
- [2] Kolassa, S., 2023. Do we want coherent hierarchical forecasts, or minimal MAPEs or MAEs?(We won't get both!). *International Journal of Forecasting*, 39(4), pp.1512-1517.
- [3] Hewamalage, Hansika, Christoph Bergmeir, and Kasun Bandara. "Global models for time series forecasting: A simulation study." *Pattern Recognition* 124 (2022): 108441.
- [4] Salinas, David, et al. "DeepAR: Probabilistic forecasting with autoregressive recurrent networks." *International Journal of Forecasting* 36.3 (2020): 1181-1191.
- [5] Hyndman, Rob J., and Anne B. Koehler. "Another look at measures of forecast accuracy." *International journal of forecasting* 22.4 (2006): 679-688.

Any questions?



Slides



<https://github.com/KishManani/PyDataLondon2024>

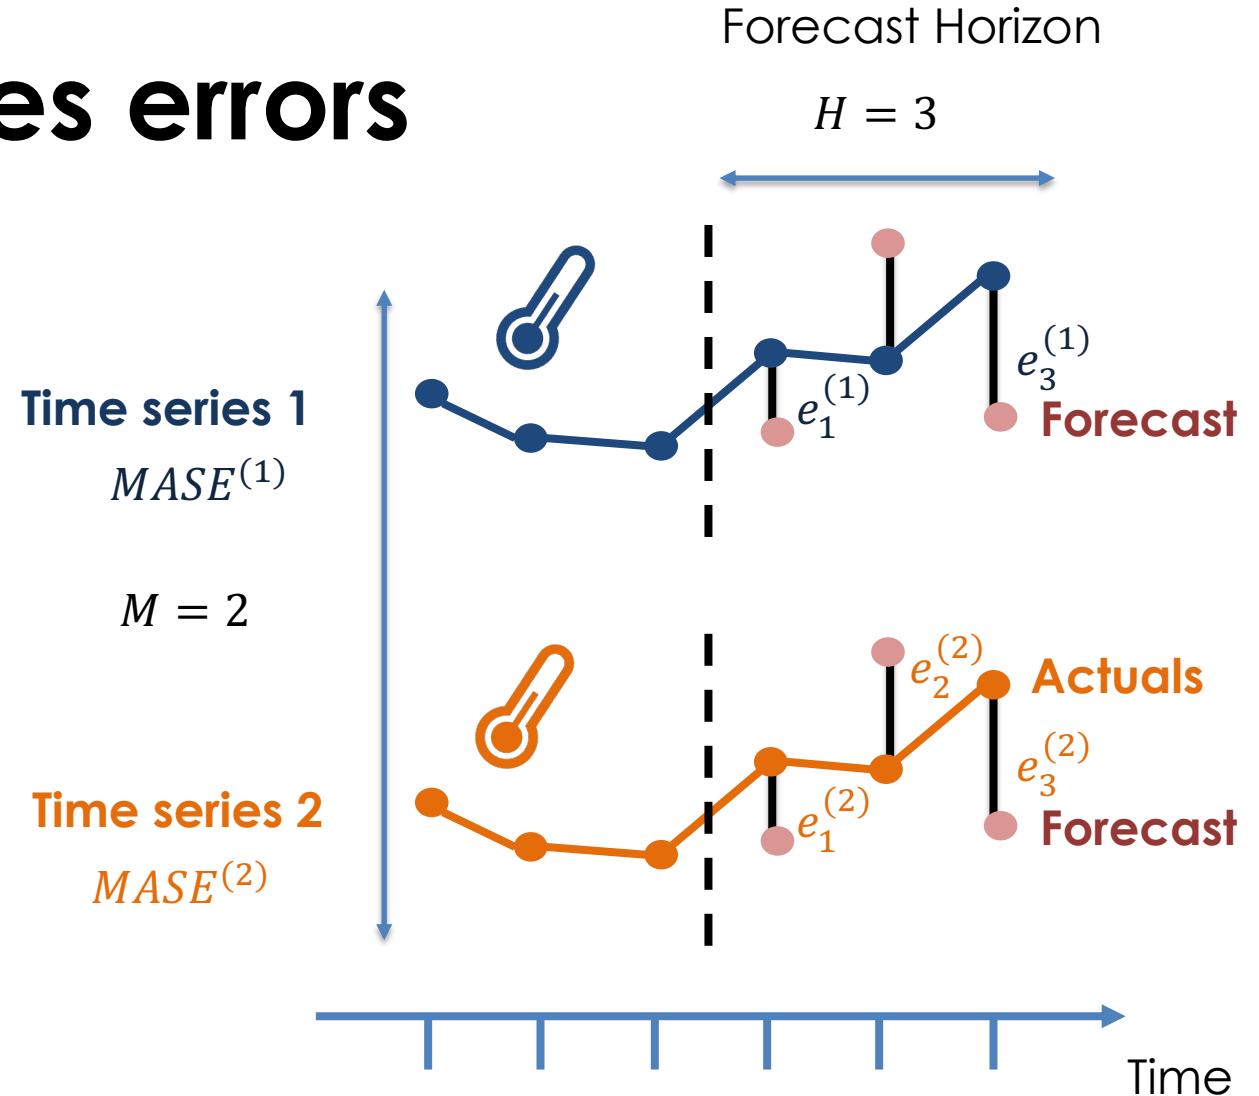
Appendix

Multiple time series errors

Average the scale-independent error metrics

$$\text{mean MASE} = \frac{\sum_{i=1}^M \text{MASE}^{(i)}}{M}$$

$$= \frac{1}{2}(\text{MASE}^{(1)} + \text{MASE}^{(2)})$$



Multiple time series error metrics

Take the **average** of a **scale independent** error metric.

$$\text{weighted mean RMSSE} = \frac{\sum_{i=1}^M w_i \text{RMSSE}^{(i)}}{\sum_{i=1}^M w_i}$$

↑
Average over each time series

$$\text{weighted mean WAPE} = \frac{\sum_{i=1}^M w_i \text{WAPE}^{(i)}}{\sum_{i=1}^M w_i}$$

↑
Average over each time series

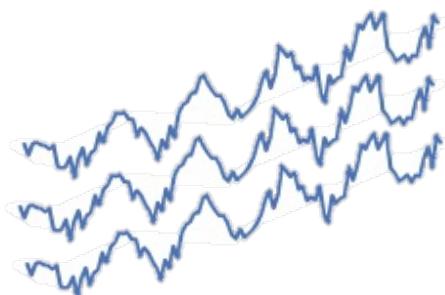
A scale independent metric is computed for each time series.

Then a weight is used to give more importance to larger scale time series.

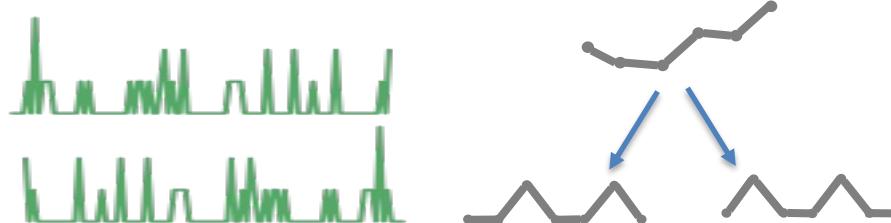
We can give equal weight if we care about each series equally.

Guidelines: Multiple time series & scale independence

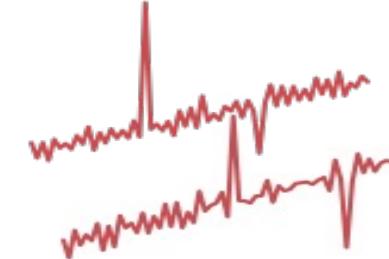
Well-behaved



Intermittent and/or hierarchical



Outliers



Average of a scale-independent error

RMSSE, MASE, WAPE

Squared errors

- RMSSE

Absolute errors & percentage errors

- WAPE
- MASE

Absolute errors

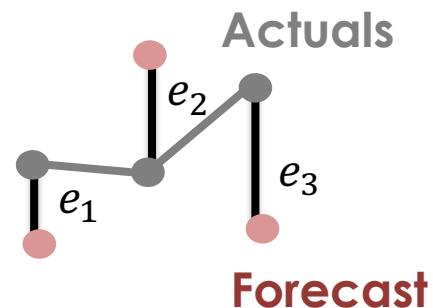
- WAPE
- MASE

Squared errors

- RMSSE

Base errors

Scale-dependent
base error



$$e_t = y_t - \hat{y}_t$$

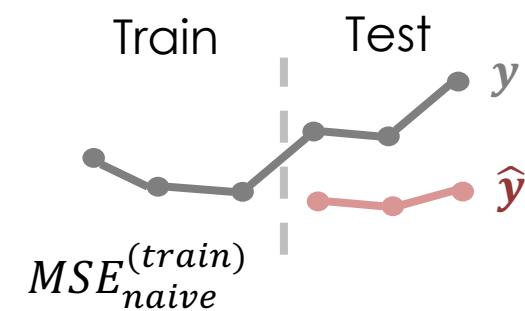
Percentage error



$$p_t$$

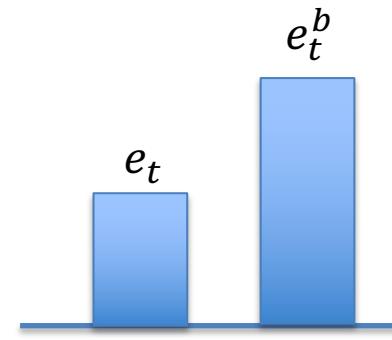
Scale-independent
base errors

Scaled error



$$q_t$$

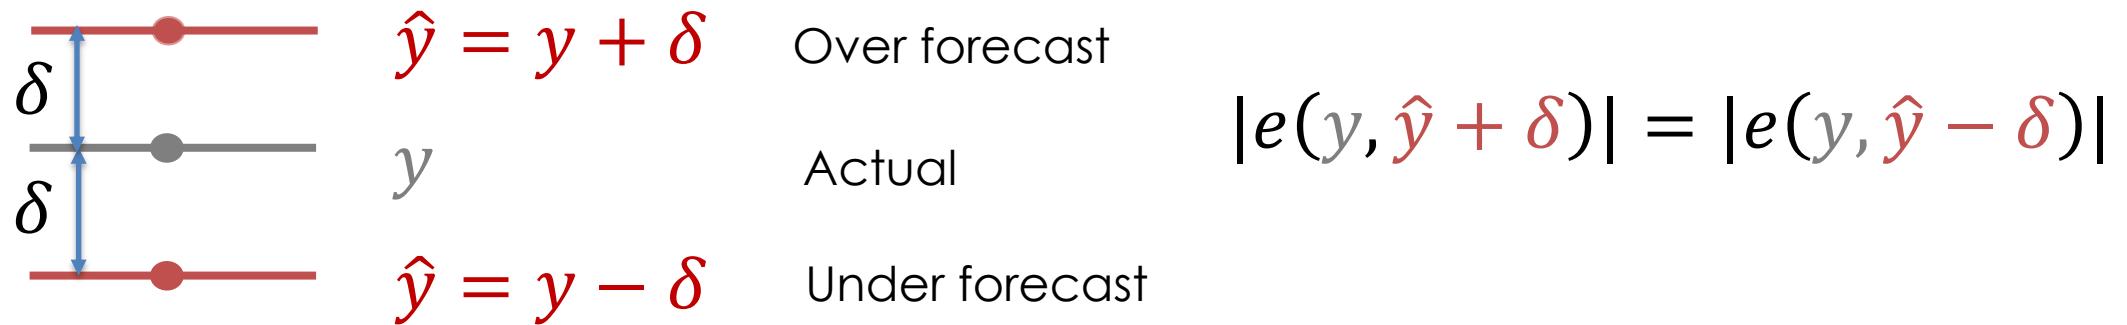
Relative error



Base errors: scale-dependent base error

- Symmetric to over/under forecasting when $\hat{y} = y \pm \delta$.

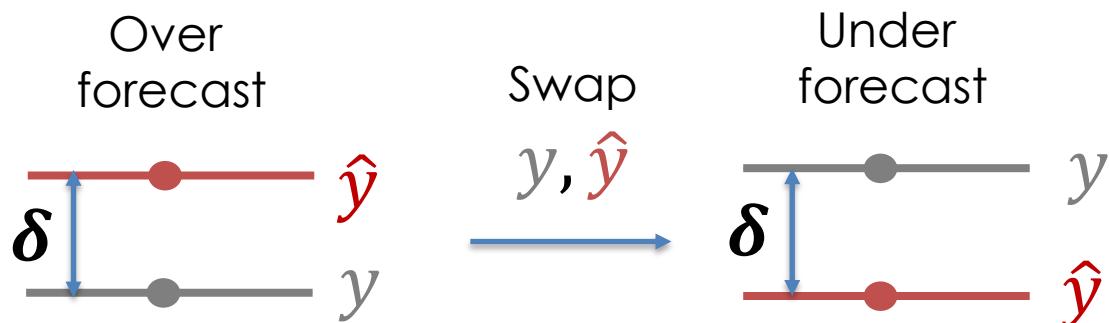
$$e_t = y_t - \hat{y}_t$$



Base errors: scale-dependent base error

- Symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✓
- Symmetric to over/under forecasting when $\hat{y} \leftrightarrow y$. ✓

$$e_t = y_t - \hat{y}_t$$



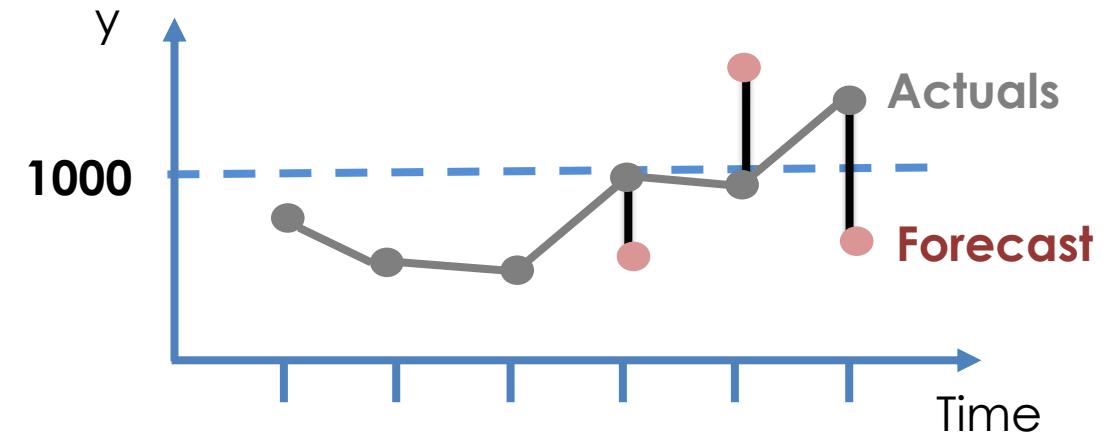
$$|e(y, \hat{y})| = |e(\hat{y}, y)|$$

Base errors: scale-dependent base error

- Symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✓
- Symmetric to over/under forecasting when $\hat{y} \leftrightarrow y$. ✓
- No issues when $y_t = 0$ or $\hat{y}_t = 0$. ✓
- Scale-dependent. !

$$e_t = y_t - \hat{y}_t$$

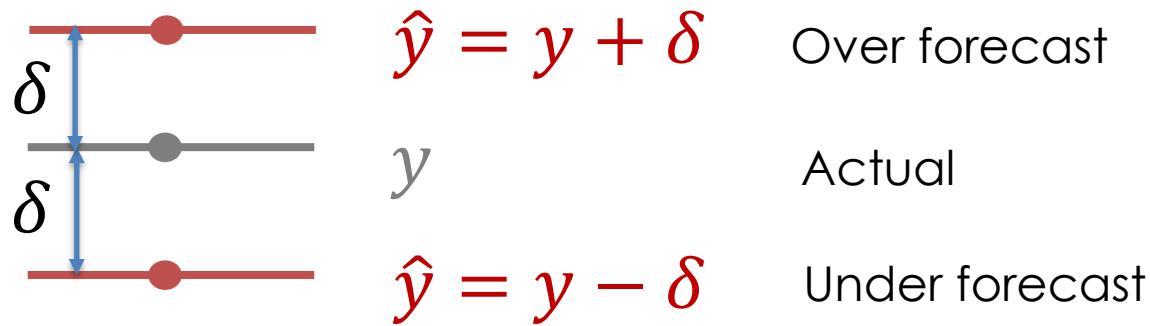
MAE = 100



Base errors: percentage error

- Scale independent. ✓
- Interpretable. ✓
- Symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✓

$$p_t = \frac{100e_t}{y_t} = 100 \frac{y_t - \hat{y}_t}{y_t}$$

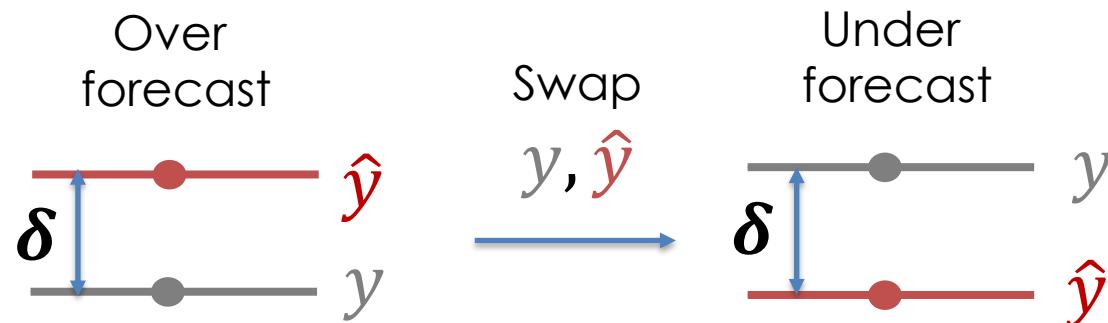


$$e(y, \hat{y} + \delta) = e(y, \hat{y} - \delta)$$

Base errors: percentage error

- Scale independent. ✓
- Interpretable. ✓
- Symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✓
- Not symmetric to over/under forecasting when $\hat{y} \leftrightarrow y$. ✗

$$p_t = \frac{100e_t}{y_t} = 100 \frac{y_t - \hat{y}_t}{y_t}$$



$$e(y, \hat{y}) \neq e(\hat{y}, y)$$

Base errors: percentage error

- Scale independent. ✓
- Interpretable. ✓
- Symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✓
- Not symmetric to over/under forecasting when $\hat{y} \leftrightarrow y$. ✗
- Unstable when y_t is small. ✗

$$p_t = \frac{100e_t}{y_t} = 100 \frac{y_t - \hat{y}_t}{y_t}$$

Base errors: percentage error

- Scale independent. ✓
- Interpretable. ✓
- Symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✓
- Not symmetric to over/under forecasting when $\hat{y} \leftrightarrow y$. ✗
- Unstable when y_t is small. ✗
- When $y_t = 0$ p_t is undefined. ✗

$$p_t = \frac{100e_t}{y_t} = 100 \frac{y_t - \hat{y}_t}{y_t}$$

$$p_t = 100 \frac{0 - \hat{y}_t}{0} = \infty; \quad y_t = 0$$

Base errors: percentage error

- Scale independent. ✓
- Interpretable. ✓
- Symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✓
- Not symmetric to over/under forecasting when $\hat{y} \leftrightarrow y$. ✗
- Unstable when y_t is small. ✗
- When $y_t = 0$ p_t is undefined. ✗
- When $\hat{y}_t = 0$ p_t is always 100%, regardless of e_t . ✗

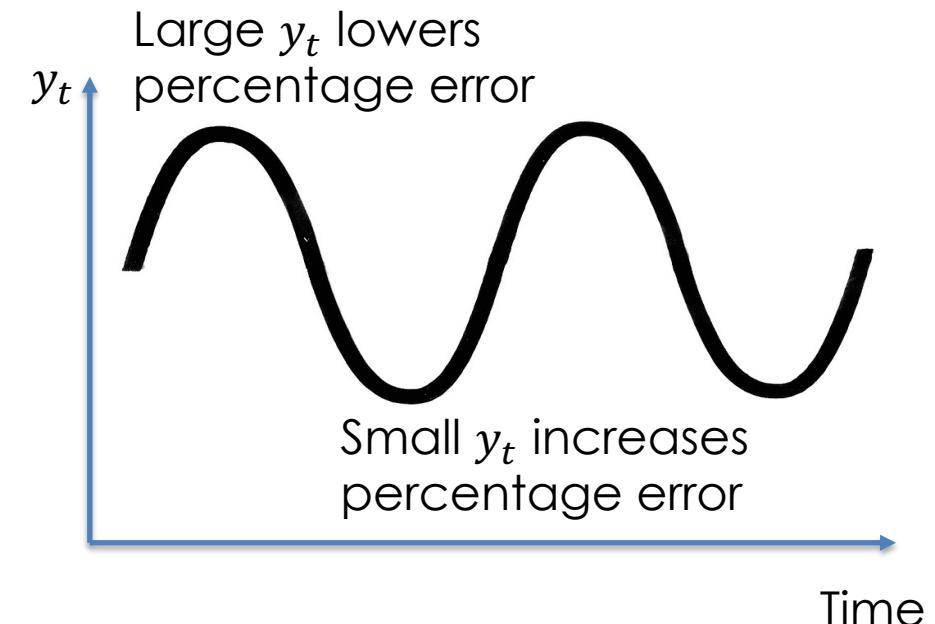
$$p_t = \frac{100e_t}{y_t} = 100 \frac{y_t - \hat{y}_t}{y_t}$$

$$p_t = 100 \frac{y_t - 0}{y_t} = 100; \hat{y}_t = 0$$

Base errors: percentage error

- Scale independent. ✓
- Interpretable. ✓
- Symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✓
- Not symmetric to over/under forecasting when $\hat{y} \leftrightarrow y$. ✗
- Unstable when y_t is small. ✗
- When $y_t = 0$ p_t is undefined. ✗
- When $\hat{y}_t = 0$ p_t is always 100%, regardless of e_t . ✗
- Seasonality and outliers can inflate and deflate p_t . ✗

$$p_t = \frac{100e_t}{y_t} = 100 \frac{y_t - \hat{y}_t}{y_t}$$



Base errors: symmetric percentage error

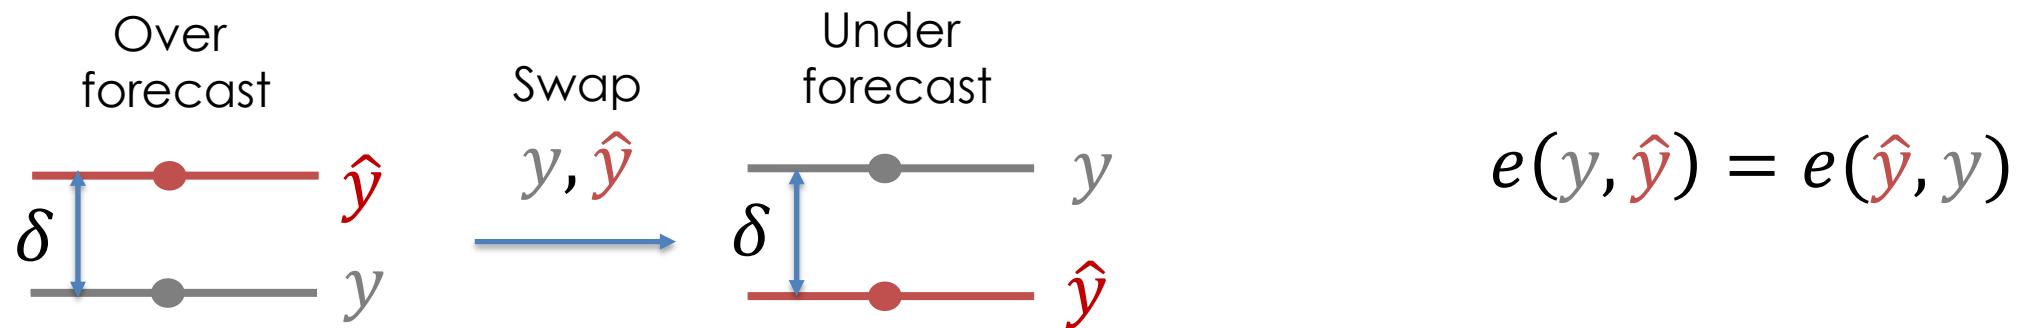
- Scale independent. 
- Less interpretable. 

$$p_t^* = \frac{100e_t}{\frac{1}{2}(|y_t| + |\hat{y}_t|)}$$

Base errors: symmetric percentage error

- Scale independent. ✓
- Less interpretable. ✗
- Symmetric to over/under forecasting when $\hat{y} \leftrightarrow y$. ✓

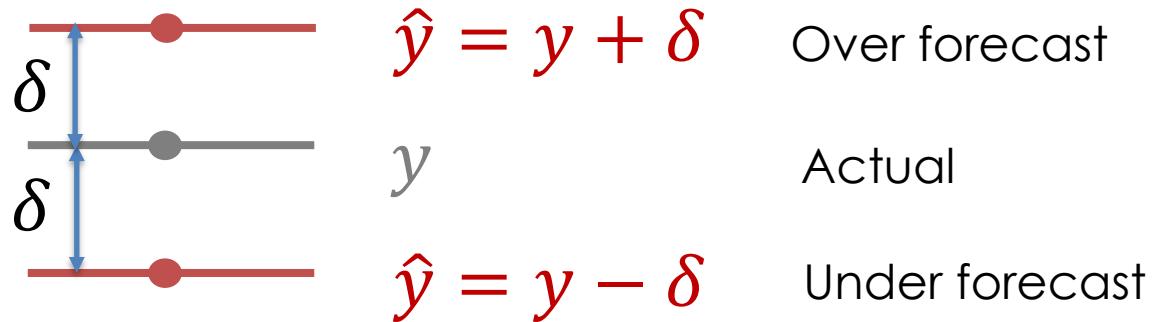
$$p_t^* = \frac{100e_t}{\frac{1}{2}(|y_t| + |\hat{y}_t|)}$$



Base errors: symmetric percentage error

- Scale independent. ✓
- Less interpretable. ✗
- Symmetric to over/under forecasting when $\hat{y} \leftrightarrow y$. ✓
- Not symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✗

$$p_t^* = \frac{100e_t}{\frac{1}{2}(|y_t| + |\hat{y}_t|)}$$



$$e(y, \hat{y} + \delta) \neq e(y, \hat{y} - \delta)$$

Base errors: symmetric percentage error

- Scale independent. ✓
- Less interpretable. ✗
- Symmetric to over/under forecasting when $\hat{y} \leftrightarrow y$. ✓
- Not symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✗
- If either $y_t = 0$ or $\hat{y}_t = 0$ then p_t^* is 200%, regardless of e_t . ✗

$$p_t^* = \frac{100e_t}{\frac{1}{2}(|y_t| + |\hat{y}_t|)}$$

$$p_t^* = 200; \hat{y}_t = 0 \text{ or } y_t = 0$$

Exercise caution or avoid using MAPE and sMAPE

Mean absolute percentage error

$$MAPE = \text{mean}(|p_t|)$$

Symmetric mean absolute percentage error

$$sMAPE = \text{mean}(|p_t^*|)$$

WAPE as an alternative

- Weighted mean absolute percentage error (WAPE) can be helpful overcoming weaknesses.
- Scale independent when comparing series. ✓
- Interpretable. ✓
- Symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✓
- Symmetric to over/under forecasting when $\hat{y} \leftrightarrow y$. ✓
- Well defined when $\hat{y}_t = 0$. ✓
- Well defined when some $y_t = 0$ in horizon. ✓
- Not well defined when all $y_t = 0$ in horizon. ✗
- Outliers can inflate/deflate the error. ✗

Weighted mean absolute percentage error

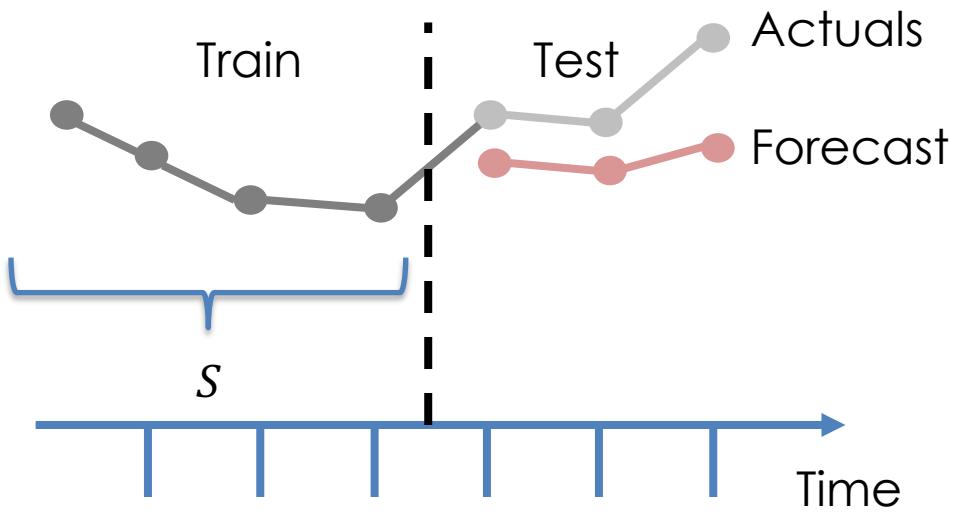
$$\begin{aligned} WAPE &= \frac{1}{\sum_t w_t} \sum_{t=1}^H w_t |p_t| \\ \xrightarrow{w_t=|y_t|} &= \frac{1}{\sum_t |y_t|} \sum_{t=1}^H |e_t| \\ &= \frac{MAE}{mean(|y_t|)} \end{aligned}$$

Base errors: scaled error

- Introduced in [Hyndman & Koehler 2005](#)

$$q_t = \frac{e_t}{S}$$

MAE or MSE of **1-step or seasonal naïve forecast on the training set**



$$S = MAE_{naive} = \frac{1}{T-1} \sum_{t=2}^T |y_t - y_{t-1}|$$

$$S = MSE_{naive} = \frac{1}{T-1} \sum_{t=2}^T (y_t - y_{t-1})^2$$

Base errors: scaled error

- Scale independent. ✓
- Interpretable. ?
- Symmetric to over/under forecasting when $\hat{y} = y \pm \delta$. ✓
- Symmetric to under/over forecasting when $\hat{y} \leftrightarrow y$. ✓.
- Well defined when $\hat{y}_t = 0$ or $y_t = 0$. ✓
- Not well defined when y_t is constant in train set. ✗
- Scale of the errors in the training set may not be the same in the future. ✗

$$q_t = \frac{e_t}{S}$$

Mean absolute scaled error

$$MASE = \text{mean}(|q_t|)$$

$$S = MAE_{naive}^{(train)}$$

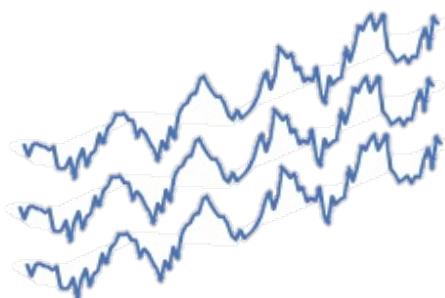
Root mean squared scaled error

$$RMSSE = \sqrt{\text{mean}(q_t^2)}$$

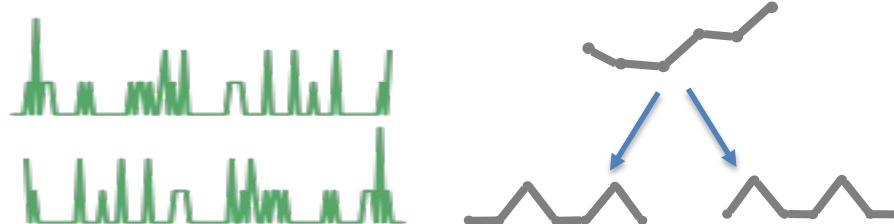
$$S = MSE_{naive}^{(train)}$$

Guidelines: Multiple time series & scale dependence

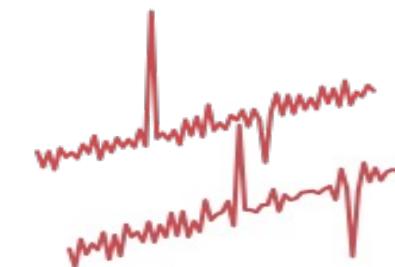
Well-behaved



Intermittent and/or hierarchical



Outliers



Pooled scale-dependent error

NRMSE, ND

Weighted average of a scale-independent error

W-RMSSE, W-MASE, W-WAPE

Squared errors

- NRMSE
- W-RMSSE

Absolute errors & percentage errors

- ND
- W-WAPE
- W-MASE

Absolute errors

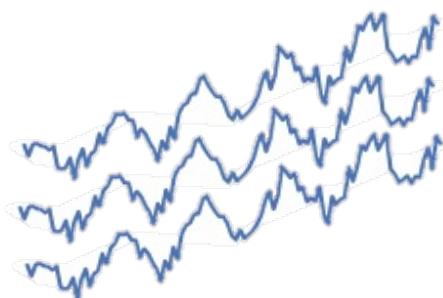
- ND
- W-MASE

Squared errors

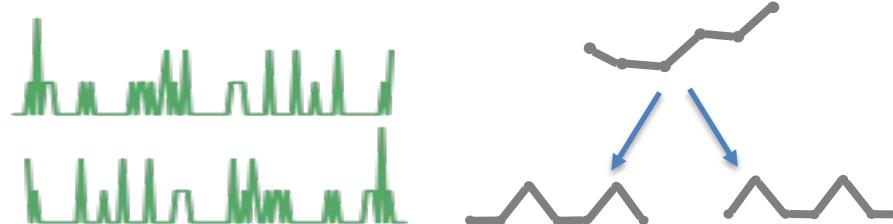
- NRMSE
- W-RMSSE

Guidelines: Multiple time series & scale independence

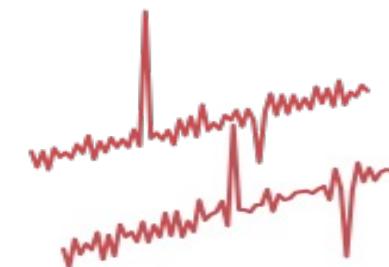
Well-behaved



Intermittent and/or hierarchical



Outliers



Average of a scale-independent error

RMSSE, MASE, WAPE

Squared errors

- RMSSE

Absolute errors & percentage errors

- WAPE
- MASE

Absolute errors

- WAPE
- MASE

Squared errors

- RMSSE

Summary operators

- **Mean**

- Sensitive to outliers.

Examples

$$MSE = \frac{1}{N} \sum_t e_t^2$$

- **Weighted mean**

- Sensitive to outliers.
- Allows us to give more importance to specific time points or values.

- **Median**

- Outlier robust.

$$WAPE = \frac{1}{\sum_t w_t} \sum_t w_t |p_t|$$

$$RMdSE = \sqrt{median(e_t^2)}$$

Transforms

- **Absolute value:** $|x|$

- Less sensitive to outliers.
- Optimises for the median.

Examples

$$MAE = \frac{1}{N} \sum_t |e_t|$$

- **Square:** x^2

- Sensitive to outliers.
- Optimises for the mean.
- Better for intermittent data.
- Better for coherent hierarchical forecasts.

$$RMSE = \sqrt{\text{mean}(e_t^2)}$$

Modified percentage errors

- Exercise caution and/or avoid using MAPE and sMAPE.
- Consider using **modified** MAPE or sMAPE to handle instability when denominator is small.
- Still asymmetric to under and over forecasting. X
- $\hat{y}_t = 0$ and $y_t = 0$ still an issue. X

Modified MAPE

$$mMAPE = \frac{1}{H} \sum_{t=1}^H \frac{100|e_t|}{|y_t| + \epsilon}$$

$\epsilon = 1$ in [Bandera et al. 2019](#)

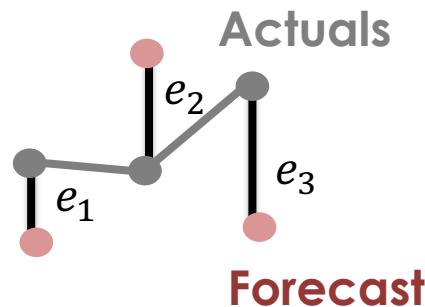
Modified sMAPE

$$msMAPE = \frac{1}{H} \sum_{t=1}^H \frac{100|e_t|}{\frac{1}{2} \max(|y_t| + |\hat{y}_t| + \epsilon, 0.5 + \epsilon)}$$

$\epsilon = 0.1$ in [Suilin 2017](#)

Base errors

Scale-dependent
base error



$$e_t = y_t - \hat{y}_t$$

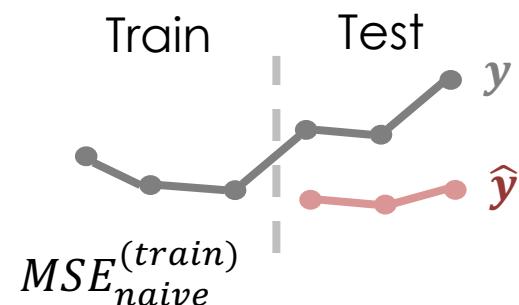
Percentage error



$$p_t = \frac{e_t}{y_t}$$

$$p_t^* = \frac{e_t}{\frac{1}{2}(|y_t| + |\hat{y}_t|)}$$

Scaled error

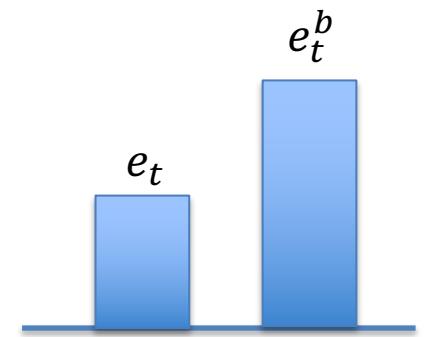


$$q_t = \frac{e_t}{S}$$

$$S = MSE_{naive}^{(train)}$$

$$S = MAE_{naive}^{(train)}$$

Relative error



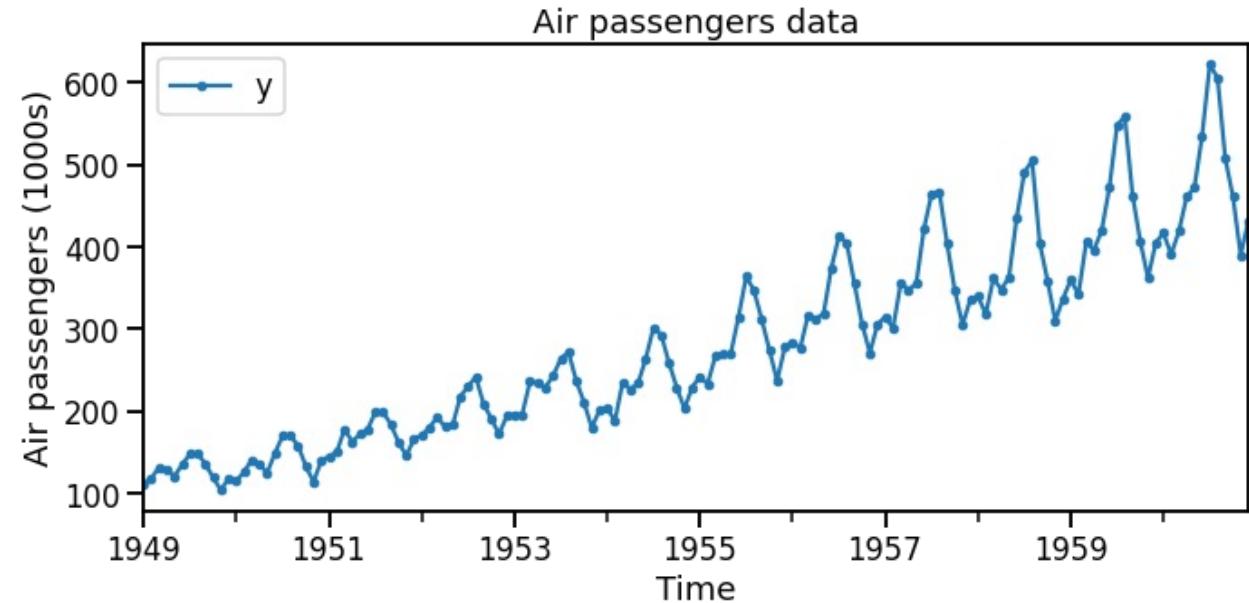
$$r_t = \frac{e_t}{e_t^b}$$

$$e_t^b = y_t - \hat{y}_t^b$$

Use simpler methods for “easy” time series

Time series characteristics

- Strong seasonality and/or trend.
- Small number of time series.
- Uncorrelated time series.
- No sparsity or intermittency.
- Few or no exogenous features.

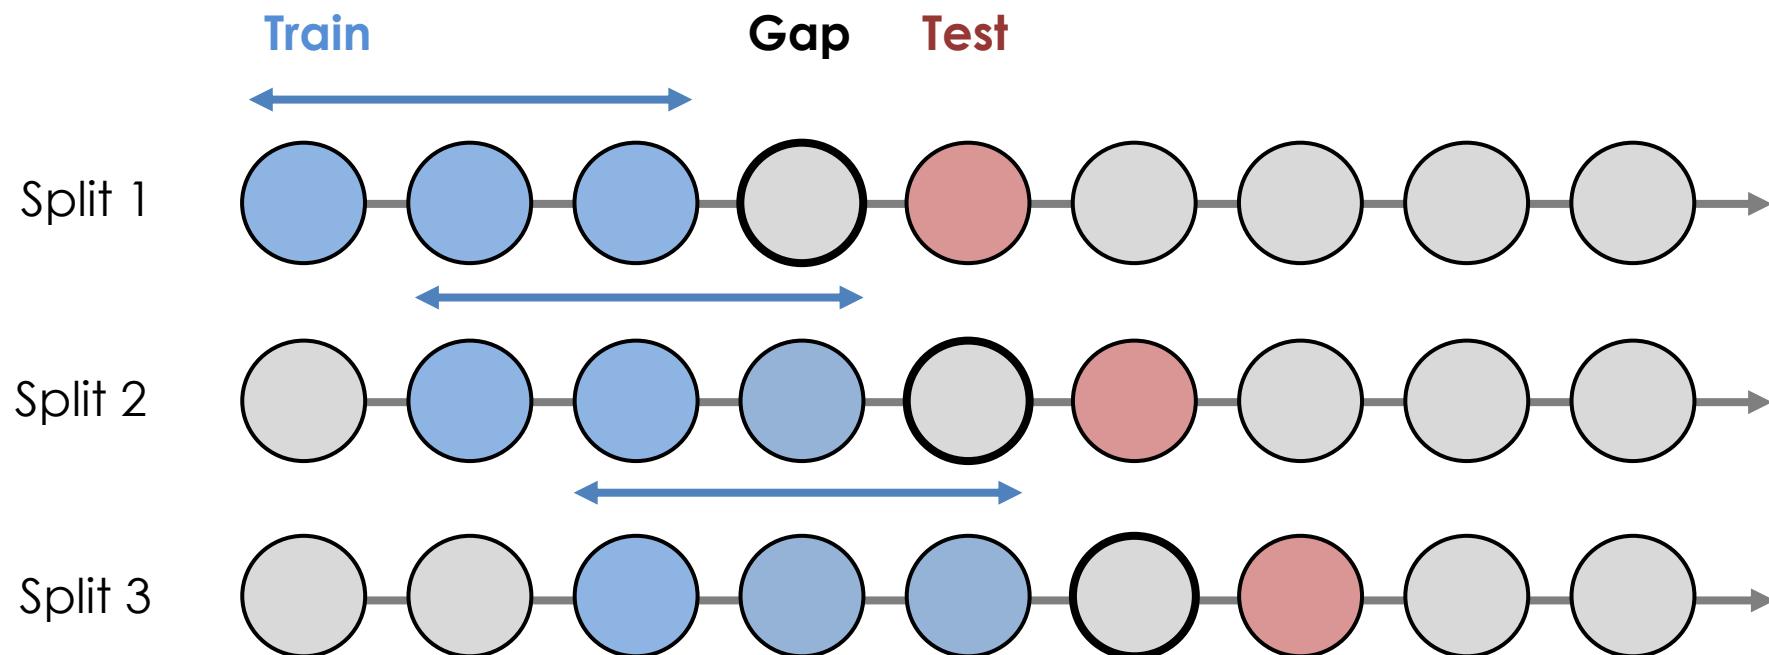


Methods

- Naïve and rules based forecasts
- ARIMA
- ETS
- Prophet

Backtesting: Try to reflect process in production

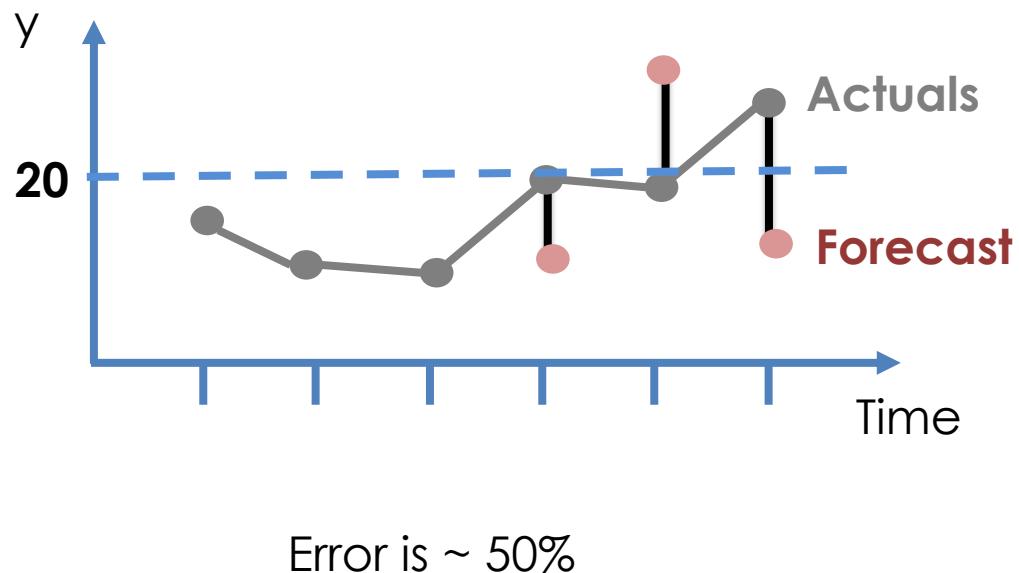
Example: There is a gap between when the forecast is created relative to when it is used. This gap should be reflected in backtesting.



Which forecast is better?

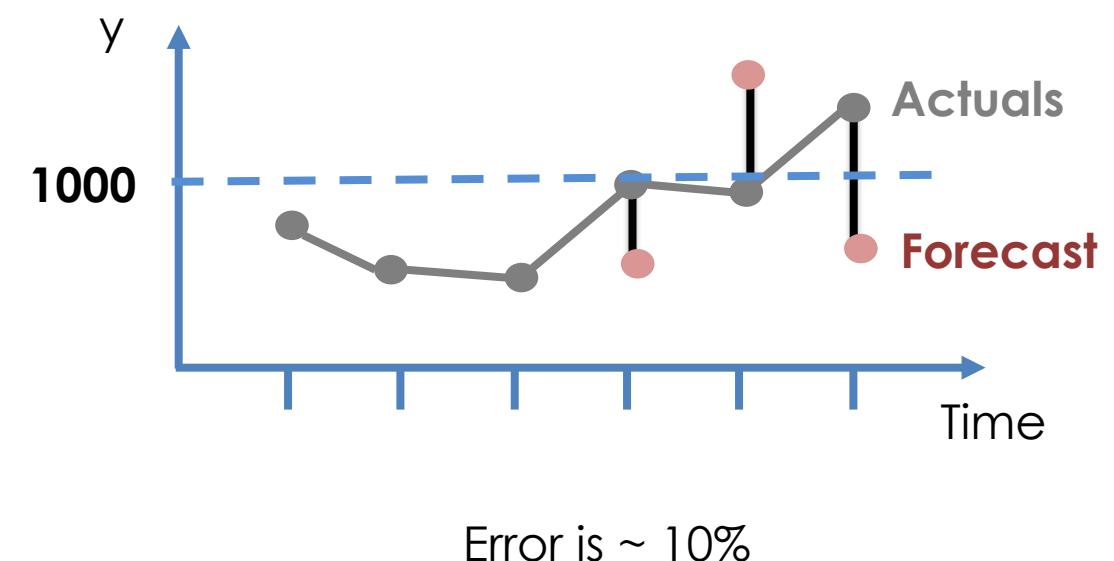
Time series 1

MAE = 10

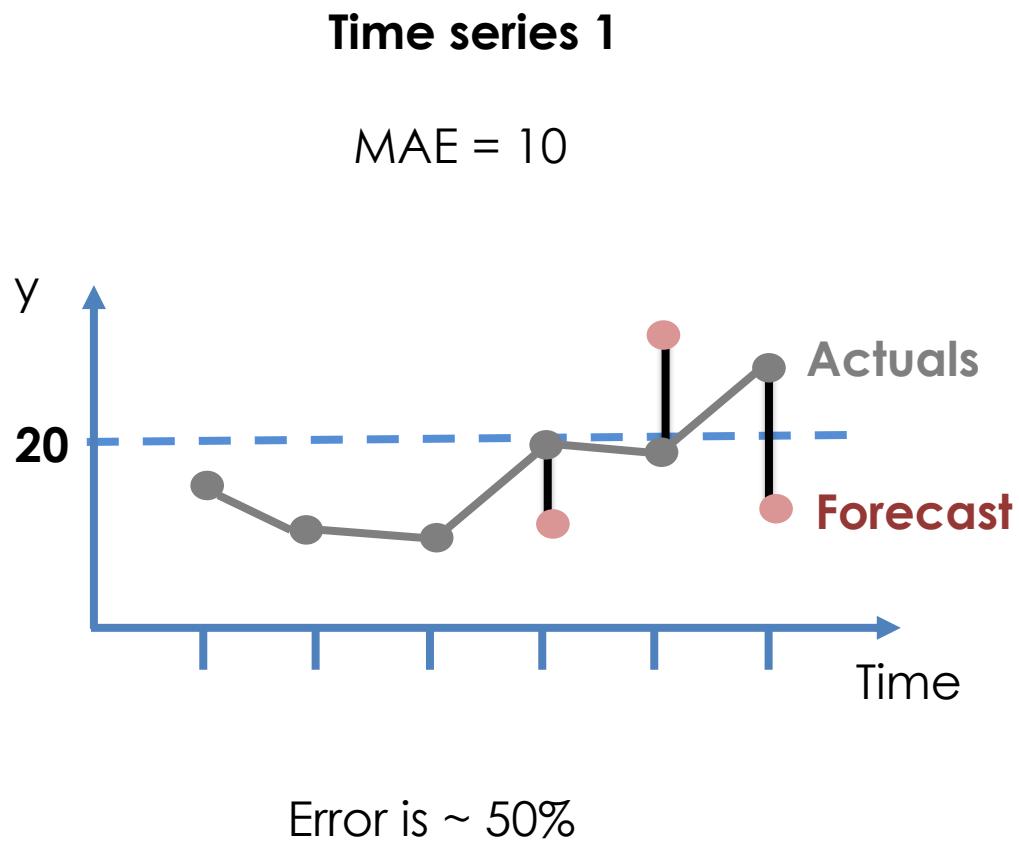


Time series 2

MAE = 100



Which forecast is better?

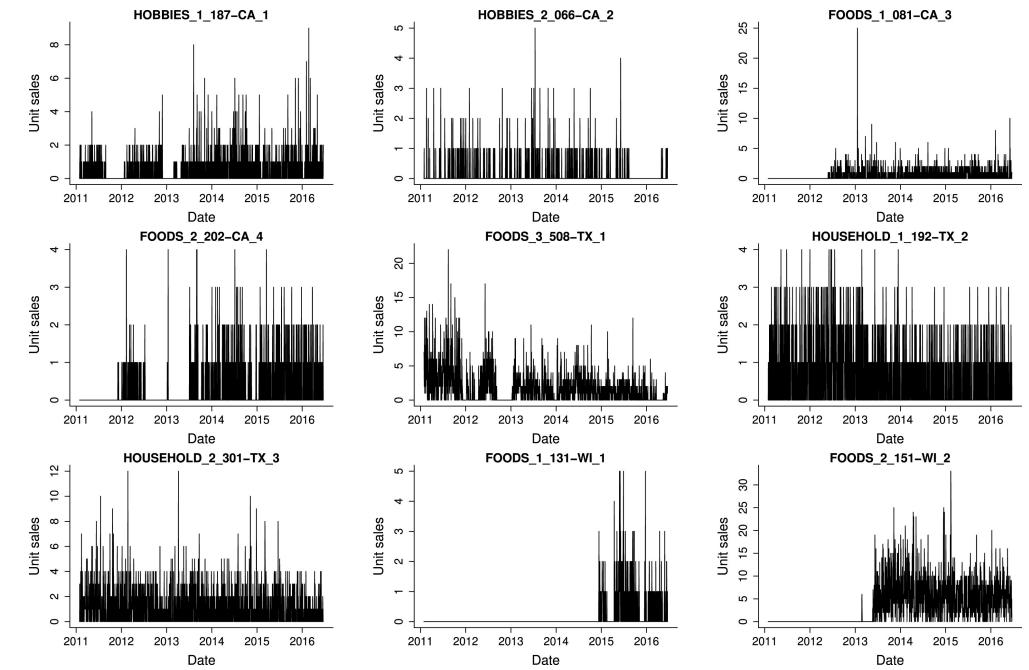
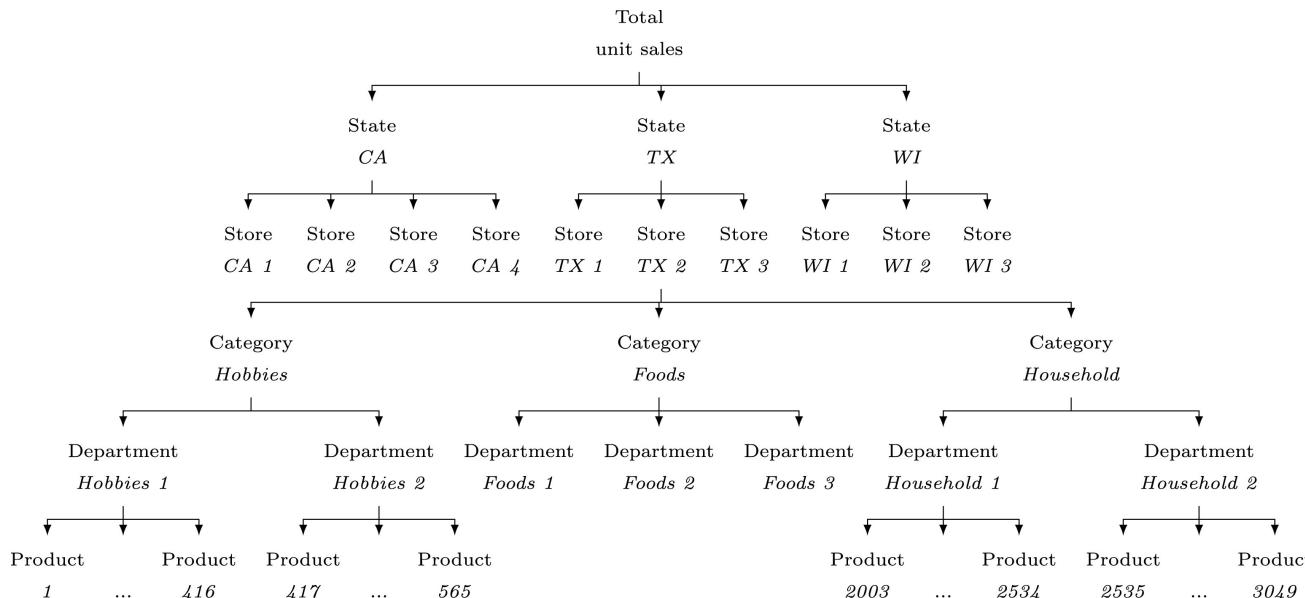


The **scale** of the time series matters.

This is important when comparing errors across **multiple time series**, across **time periods** within a series, and if there are **outliers**.

Use ML methods for “hard” time series

- Large number of correlated time series (30,490)
- Hierarchical structure
- Varying length for each time series
- High sparsity & intermittency
- Exogenous variables (price, promos, etc.)
- Multiple seasonal patterns



[1] Makridakis, Spyros, Evangelos Spiliotis, and Vassilios Assimakopoulos. "The M5 competition: Background, organization, and implementation." *International Journal of Forecasting* (2021).

Kishan Manani — in/KishanManani — trainindata.com/p/forecasting-specialization